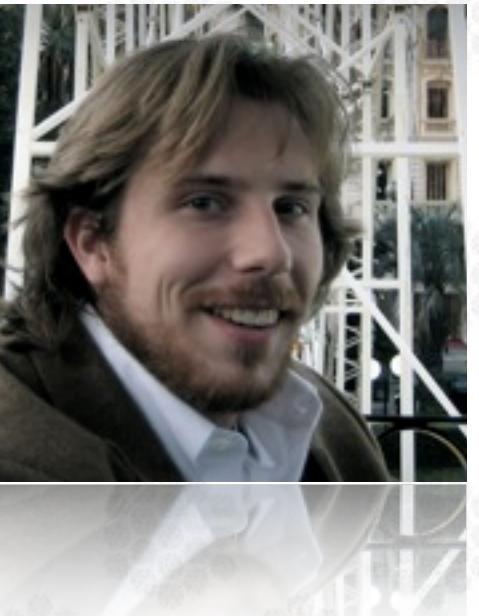


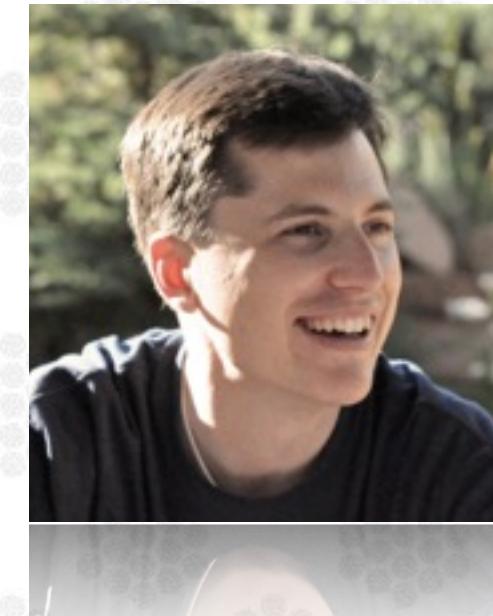
CARBON CAPTURE
MULTIDISCIPLINARY
SIMULATION CENTER

Transparent Abstractions for Scientific Computing (TASC)



James C. Sutherland

ASSOCIATE PROFESSOR, CHEMICAL ENGINEERING



Matthew Might

ASSISTANT PROFESSOR, COMPUTER SCIENCE

Tony Saad

RESEARCH ASSOCIATE - ICSE

Christoper Earl

POST-DOCTORAL RESEARCH ASSOCIATE - ICSE

Orientation

- Uintah controls & allocates resources
 - Coarse data & task parallel

Application Code

- Implements physics kernels.
- Provide interaction between user input, framework & creation of physics kernels.

Uintah-X

- Coarse data parallelism.
- Coarse task parallelism.
- Fault-tolerant runtime system.
- Scalable IO.

Orientation

- Uintah controls & allocates resources
 - Coarse data & task parallel
- When resources exceed available work, Uintah pushes resources “downward” to finer levels of parallelism
 - Fine-grained task graphs within each Uintah task
 - Fine-grained data parallel (multicore/GPU)

Application Code

- Implements physics kernels.
- Provide interaction between user input, framework & creation of physics kernels.

ExprLib

- Fine-grained task parallelism.
- Interface to various structured frameworks.
- Dynamic assembly of problem structure.

SpatialOps

- **Nebo** EDSL for computational kernels.
- Fine-grained data parallelism.

Uintah-X

- Coarse data parallelism.
- Coarse task parallelism.
- Fault-tolerant runtime system.
- Scalable IO.

Orientation

- Uintah controls & allocates resources
 - Coarse data & task parallel
- When resources exceed available work, Uintah pushes resources “downward” to finer levels of parallelism
 - Fine-grained task graphs within each Uintah task
 - Fine-grained data parallel (multicore/GPU)

Wasatch

ExprLib

SpatialOps

Uintah-X

Application Code

- Implements physics kernels.
- Provide interaction between user input, framework & creation of physics kernels.

ExprLib

- Fine-grained task parallelism.
- Interface to various structured frameworks.
- Dynamic assembly of problem structure.

SpatialOps

- **Nebo** EDSL for computational kernels.
- Fine-grained data parallelism.

Uintah-X

- Coarse data parallelism.
- Coarse task parallelism.
- Fault-tolerant runtime system.
- Scalable IO.

Orientation

- Uintah controls & allocates resources
 - Coarse data & task parallel
- When resources exceed available work, Uintah pushes resources “downward” to finer levels of parallelism
 - Fine-grained task graphs within each Uintah task
 - Fine-grained data parallel (multicore/GPU)

Wasatch

Arches

ExprLib

SpatialOps

Uintah-X

Uintah-X

Application Code

- Implements physics kernels.
- Provide interaction between user input, framework & creation of physics kernels.

ExprLib

- Fine-grained task parallelism.
- Interface to various structured frameworks.
- Dynamic assembly of problem structure.

SpatialOps

- **Nebo** EDSL for computational kernels.
- Fine-grained data parallelism.

Uintah-X

- Coarse data parallelism.
- Coarse task parallelism.
- Fault-tolerant runtime system.
- Scalable IO.

Orientation

- Uintah controls & allocates resources

- Coarse data & task parallel

- When resources exceed available work,
Uintah pushes resources “downward” to
finer levels of parallelism

- Fine-grained task graphs within each Uintah task
- Fine-grained data parallel (multicore/GPU)

Wasatch

Arches

Arches 2.0?

ExprLib

SpatialOps

Uintah-X

Uintah-X

SpatialOps

Uintah-X

Application Code

- Implements physics kernels.
- Provide interaction between user input, framework & creation of physics kernels.

ExprLib

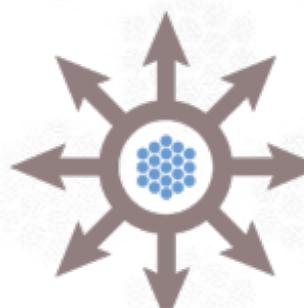
- Fine-grained task parallelism.
- Interface to various structured frameworks.
- Dynamic assembly of problem structure.

SpatialOps

- **Nebo** EDSL for computational kernels.
- Fine-grained data parallelism.

Uintah-X

- Coarse data parallelism.
- Coarse task parallelism.
- Fault-tolerant runtime system.
- Scalable IO.



Orientation

- Uintah controls & allocates resources
 - Coarse data & task parallel
- When resources exceed available work, Uintah pushes resources “downward” to finer levels of parallelism
 - Fine-grained task graphs within each Uintah task
 - Fine-grained data parallel (multicore/GPU)

Wasatch

Arches

Arches 2.0?

ODT

ExprLib

SpatialOps

Uintah-X

Uintah-X

Uintah-X

Application Code

- Implements physics kernels.
- Provide interaction between user input, framework & creation of physics kernels.

ExprLib

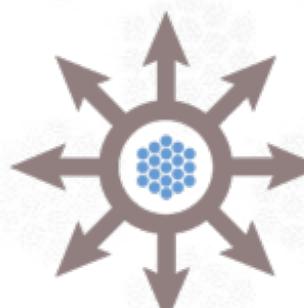
- Fine-grained task parallelism.
- Interface to various structured frameworks.
- Dynamic assembly of problem structure.

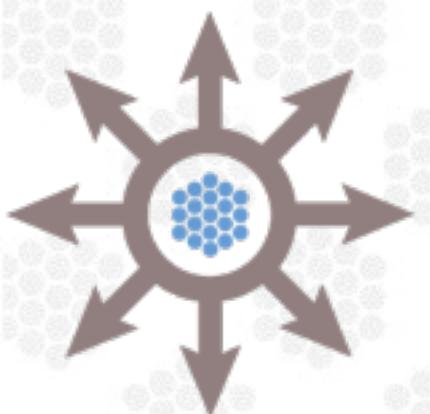
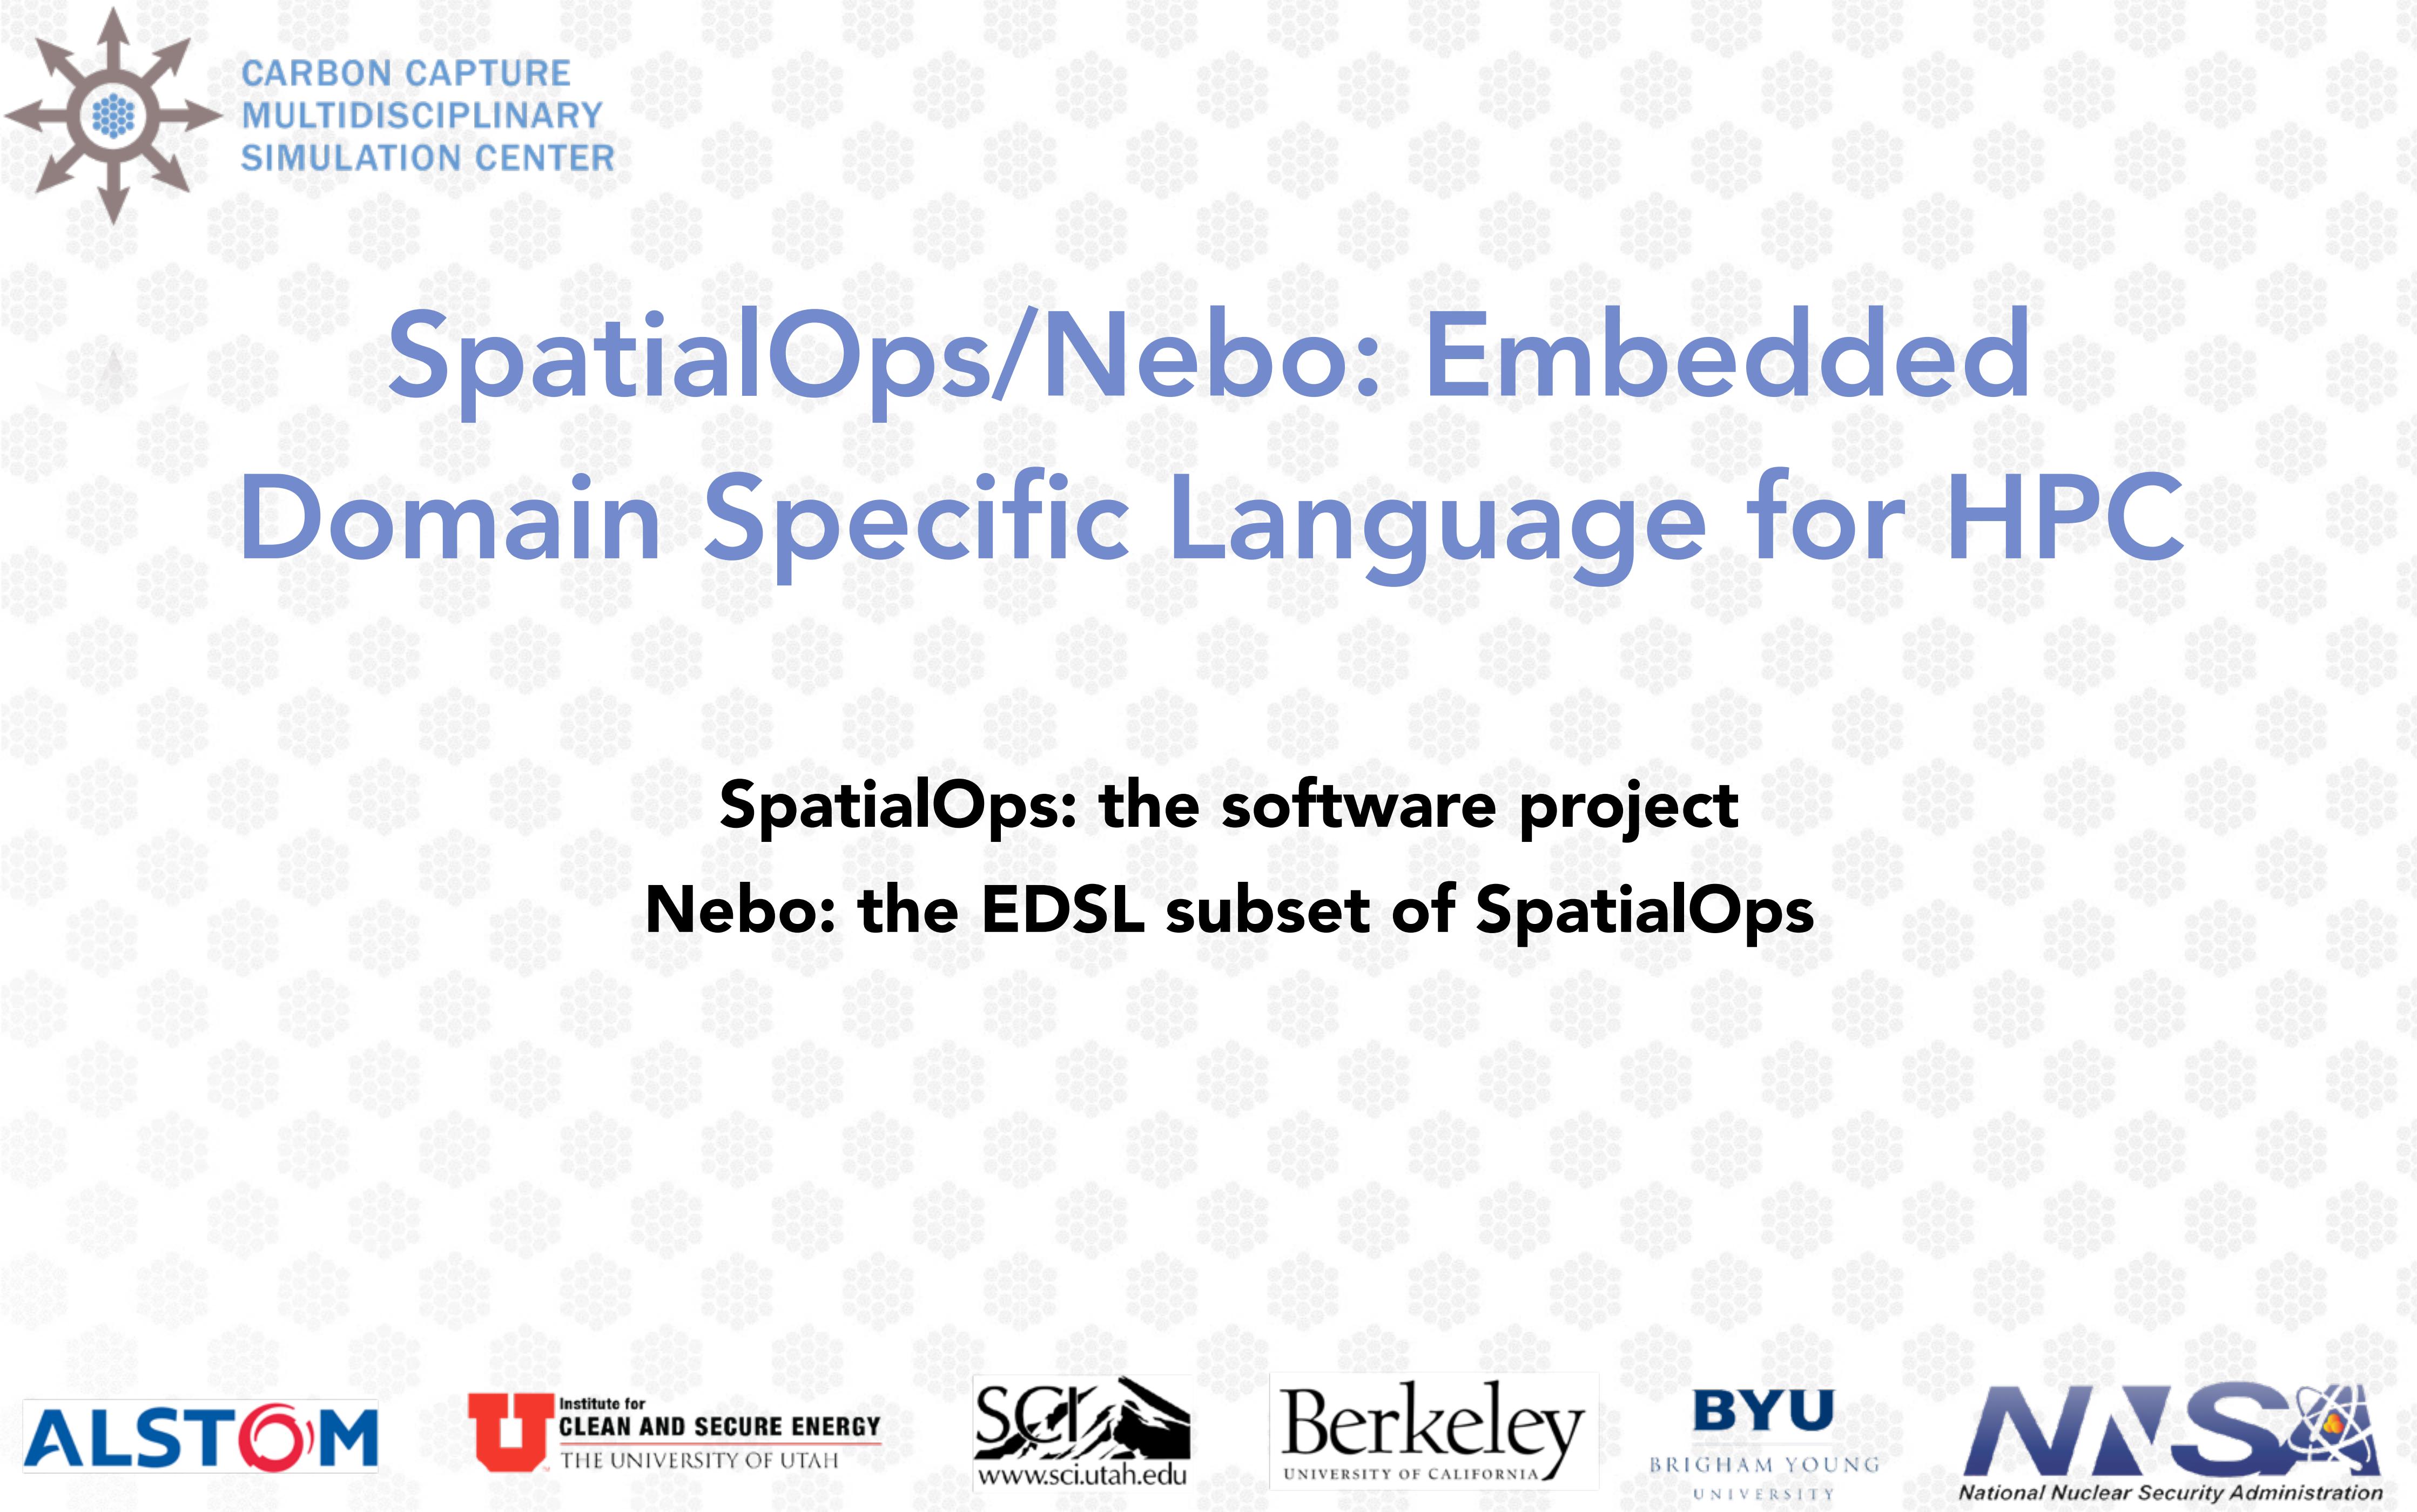
SpatialOps

- **Nebo** EDSL for computational kernels.
- Fine-grained data parallelism.

Uintah-X

- Coarse data parallelism.
- Coarse task parallelism.
- Fault-tolerant runtime system.
- Scalable IO.





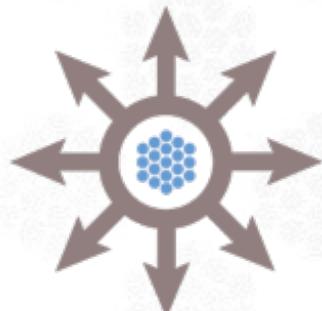
CARBON CAPTURE
MULTIDISCIPLINARY
SIMULATION CENTER

SpatialOps/Nebo: Embedded Domain Specific Language for HPC

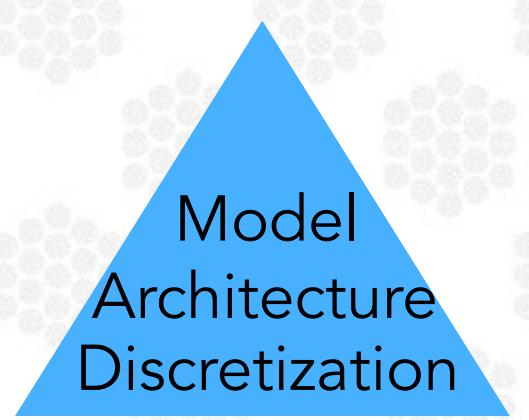
SpatialOps: the software project
Nebo: the EDSL subset of SpatialOps

Why an EDSL?

- Expressive syntax (matlab-style array operations)
 - Programmer expresses *intent* (problem structure) - not implementation.

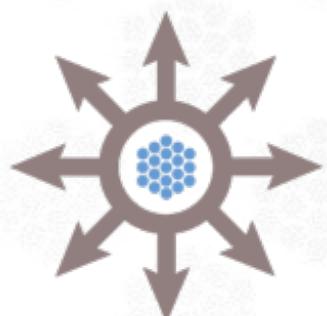
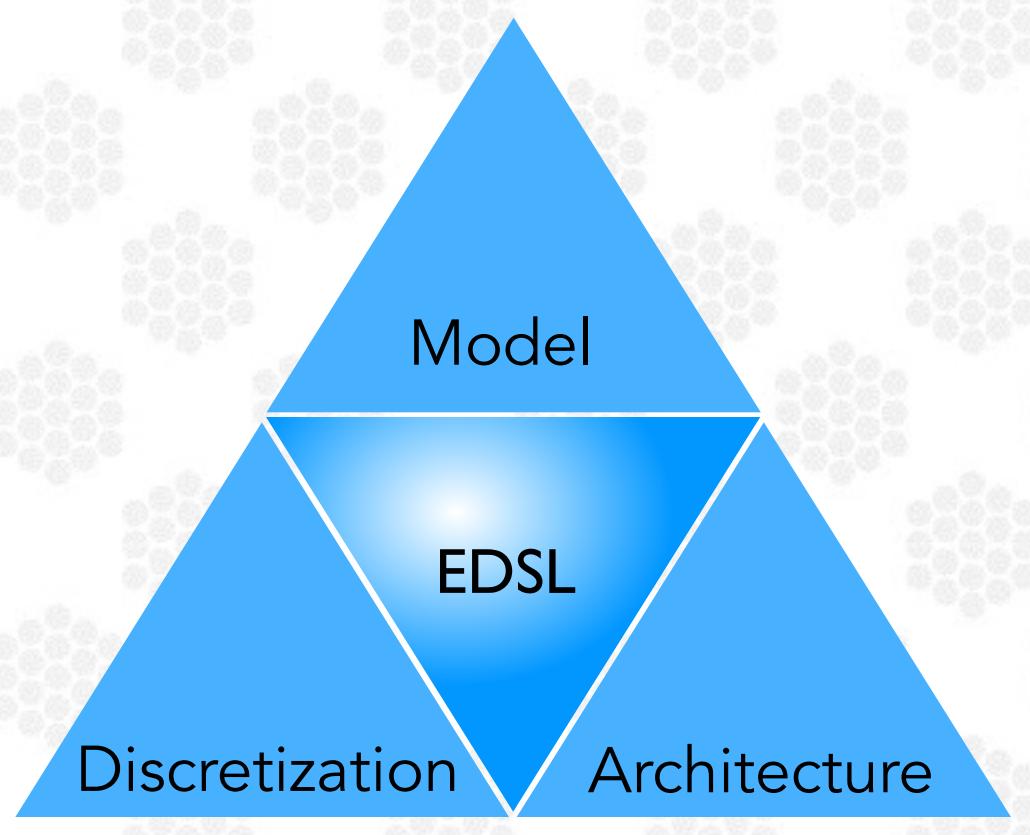


CARBON CAPTURE
MULTIDISCIPLINARY
SIMULATION CENTER



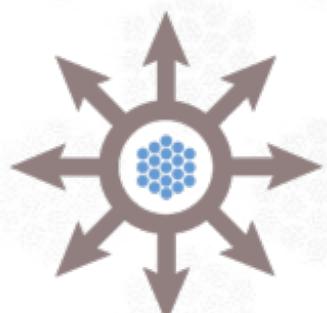
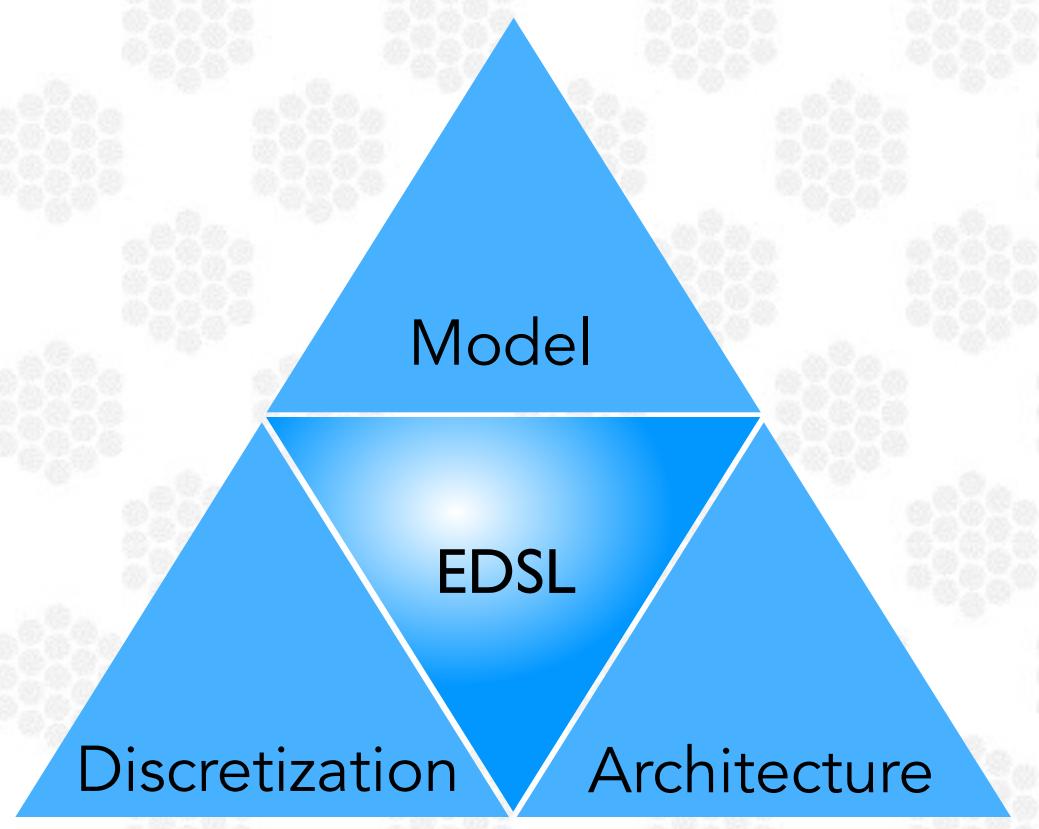
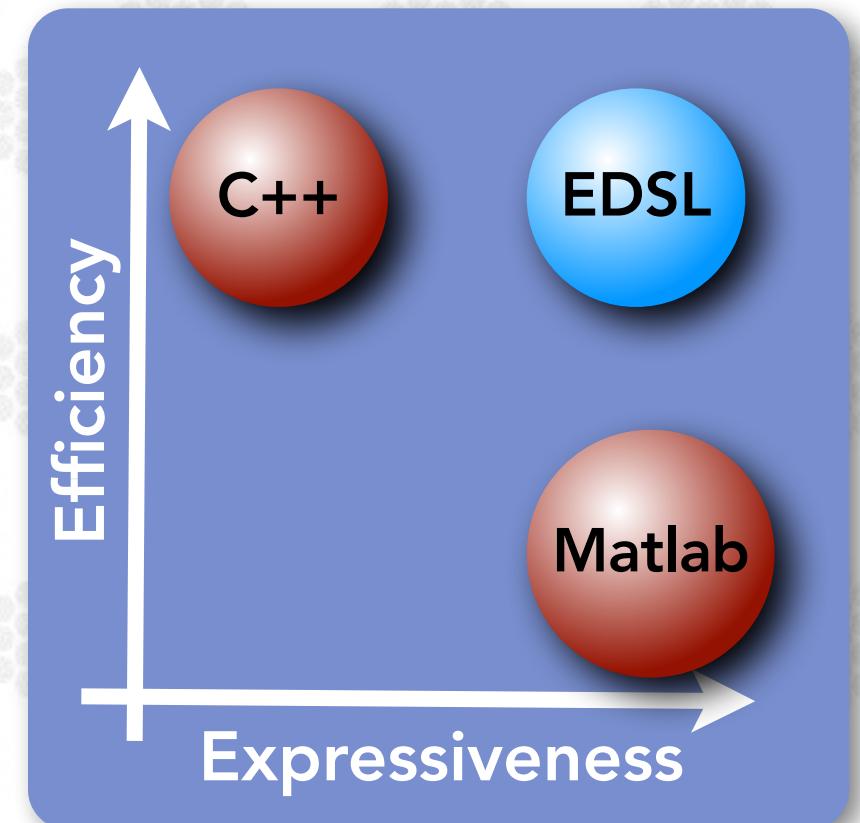
Why an EDSL?

- Expressive syntax (matlab-style array operations)
 - Programmer expresses *intent* (problem structure) - not implementation.



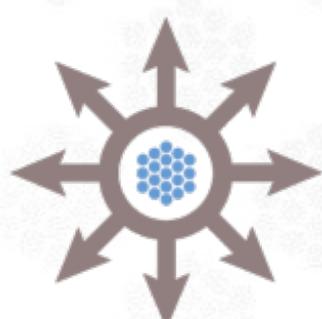
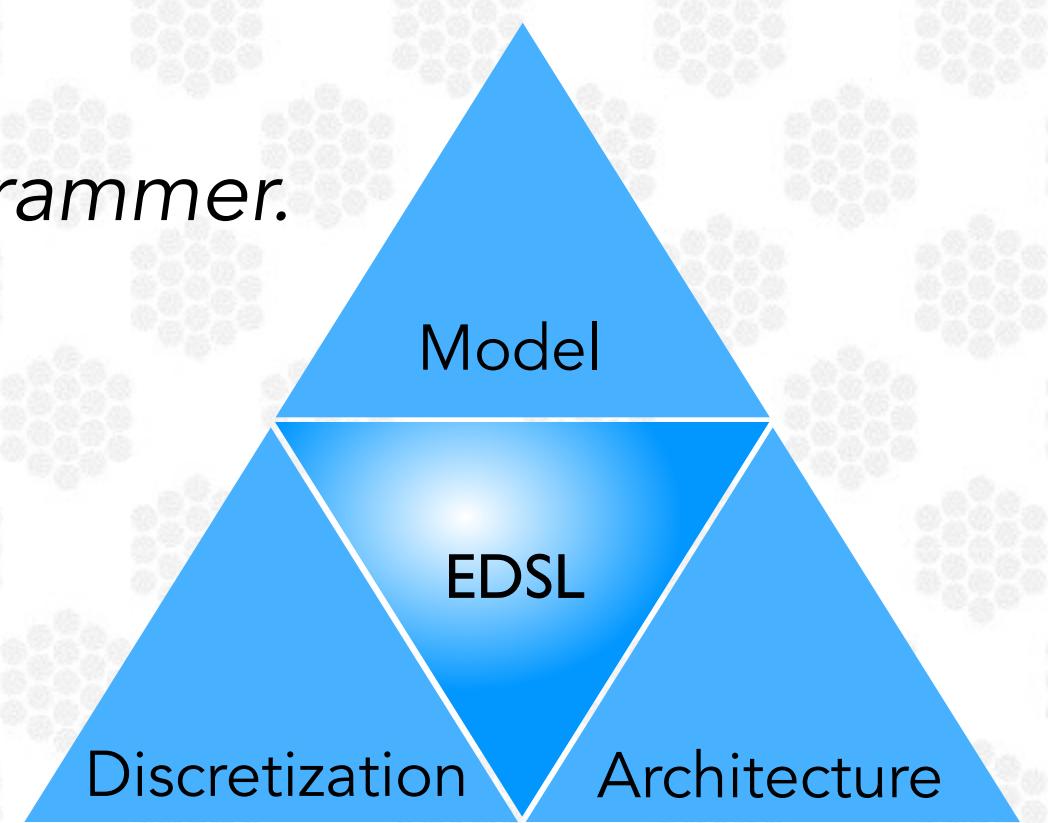
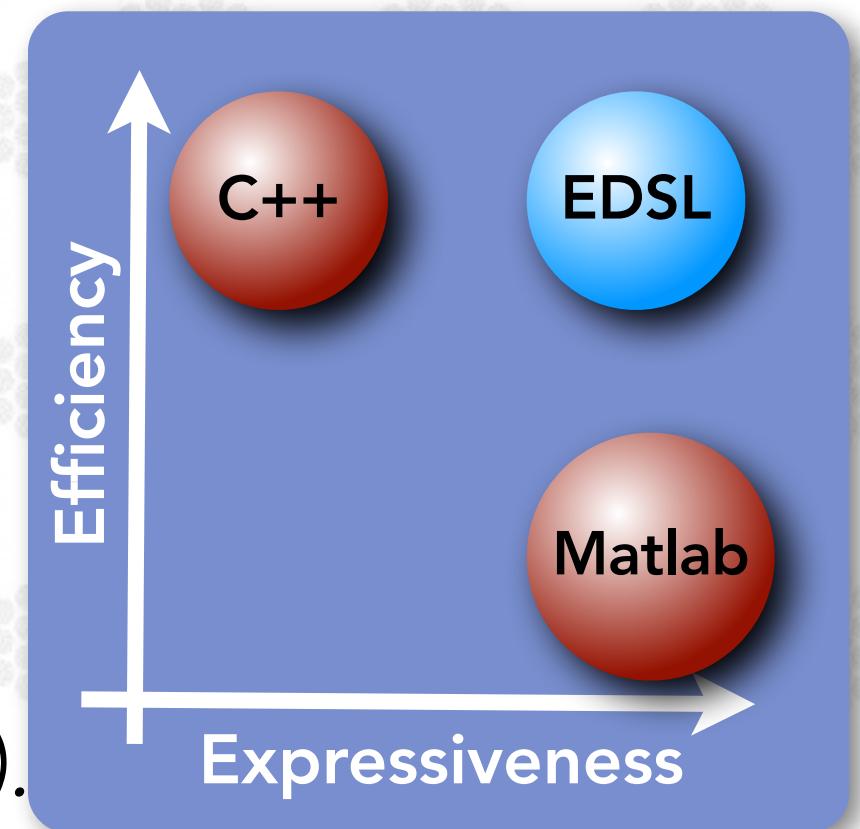
Why an EDSL?

- Expressive syntax (matlab-style array operations)
 - Programmer expresses *intent* (problem structure) - not implementation.
- High performance
 - Should match hand-tuned code in performance.



Why an EDSL?

- Expressive syntax (matlab-style array operations)
 - Programmer expresses *intent* (problem structure) - not implementation.
- High performance
 - Should match hand-tuned code in performance.
- Extensible
 - Insulate programmer from architecture changes (e.g. multicore → GPU → ...).
 - EDSL “back-end” compiles into code for target architecture.
- “Plays well with others”
 - Allow programmer to write in C++ and inter-operate with EDSL.
 - Not an “all-or-none” approach.
 - Allows concurrent development of EDSL and application codes.
 - Adapt to appropriate data structures without impacting application programmer.



Nebo Field Expressions

$$\vec{c} = \vec{a} + \sin(\vec{b})$$

Manual C++

Thread 1 {

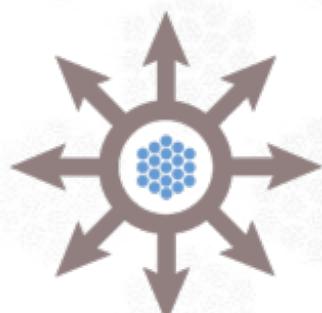
```
Field::const_iterator ia1 = a1.begin();
Field::const_iterator ib1 = b1.begin();
for(Field::iterator ic1 = c1.begin();
    ic1 != c1.end();
    ++ic1, ++ia1, ++ib1) {
    *ic1 = *ia1 + sin(*ib1);
}
...
Field::const_iterator ia_n = a_n.begin();
Field::const_iterator ib_n = b_n.begin();
for(Field::iterator ic_n = c_n.begin();
    ic_n != c_n.end();
    ++ic_n, ++ia_n, ++ib_n) {
    *ic_n = *ia_n + sin(*ib_n);
}
```

Thread n {

Nebo
→

$c <= a + \sin(b);$

- Data parallel handled internally.
- Thread deployment (resizable threadpool).
- GPU deployment.
- Compile-time guarantee of field compatibility for given operations.



Chained Stencil Operations

$$\begin{aligned}\phi &= -\nabla \cdot \mathbf{q} \\ &= \nabla \cdot (-\lambda \nabla T)\end{aligned}$$

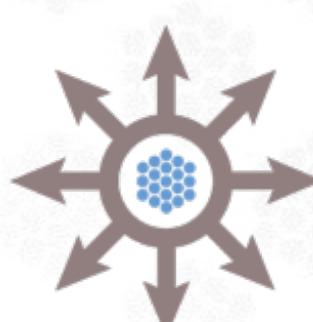
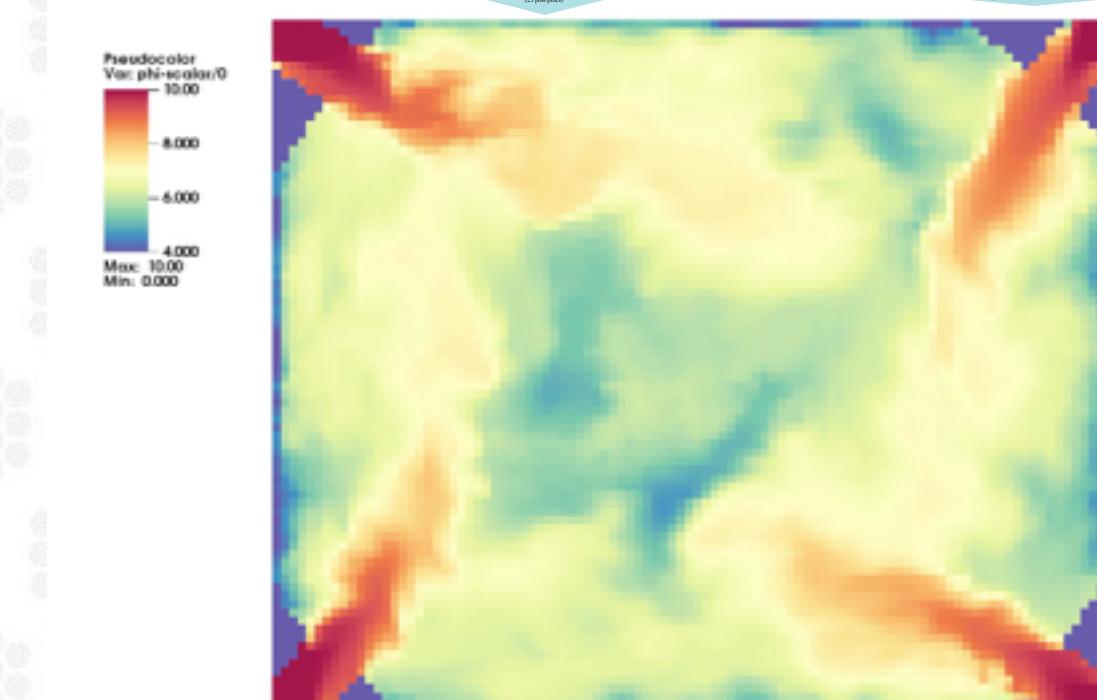
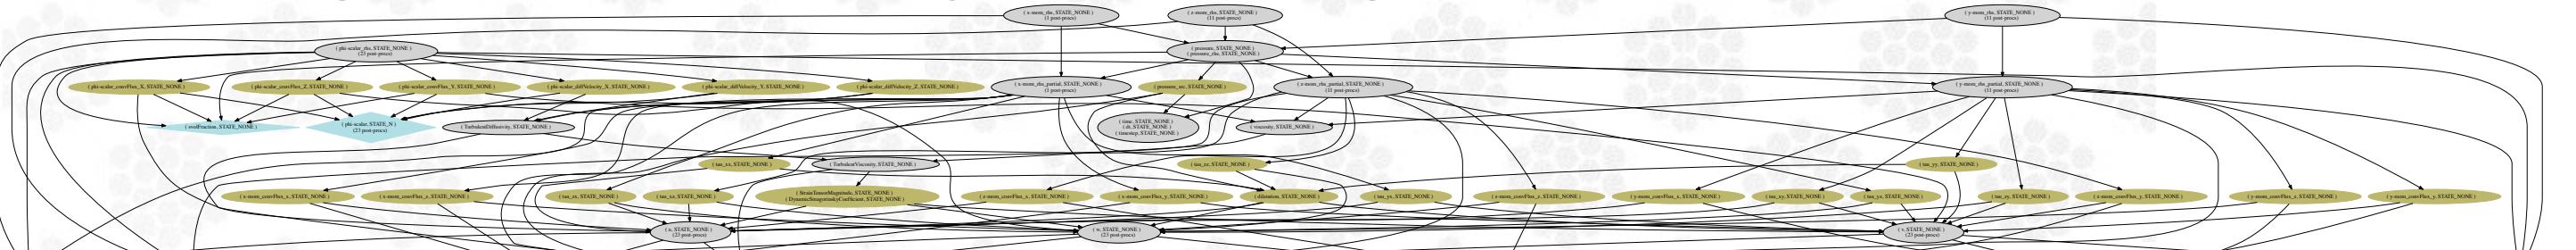
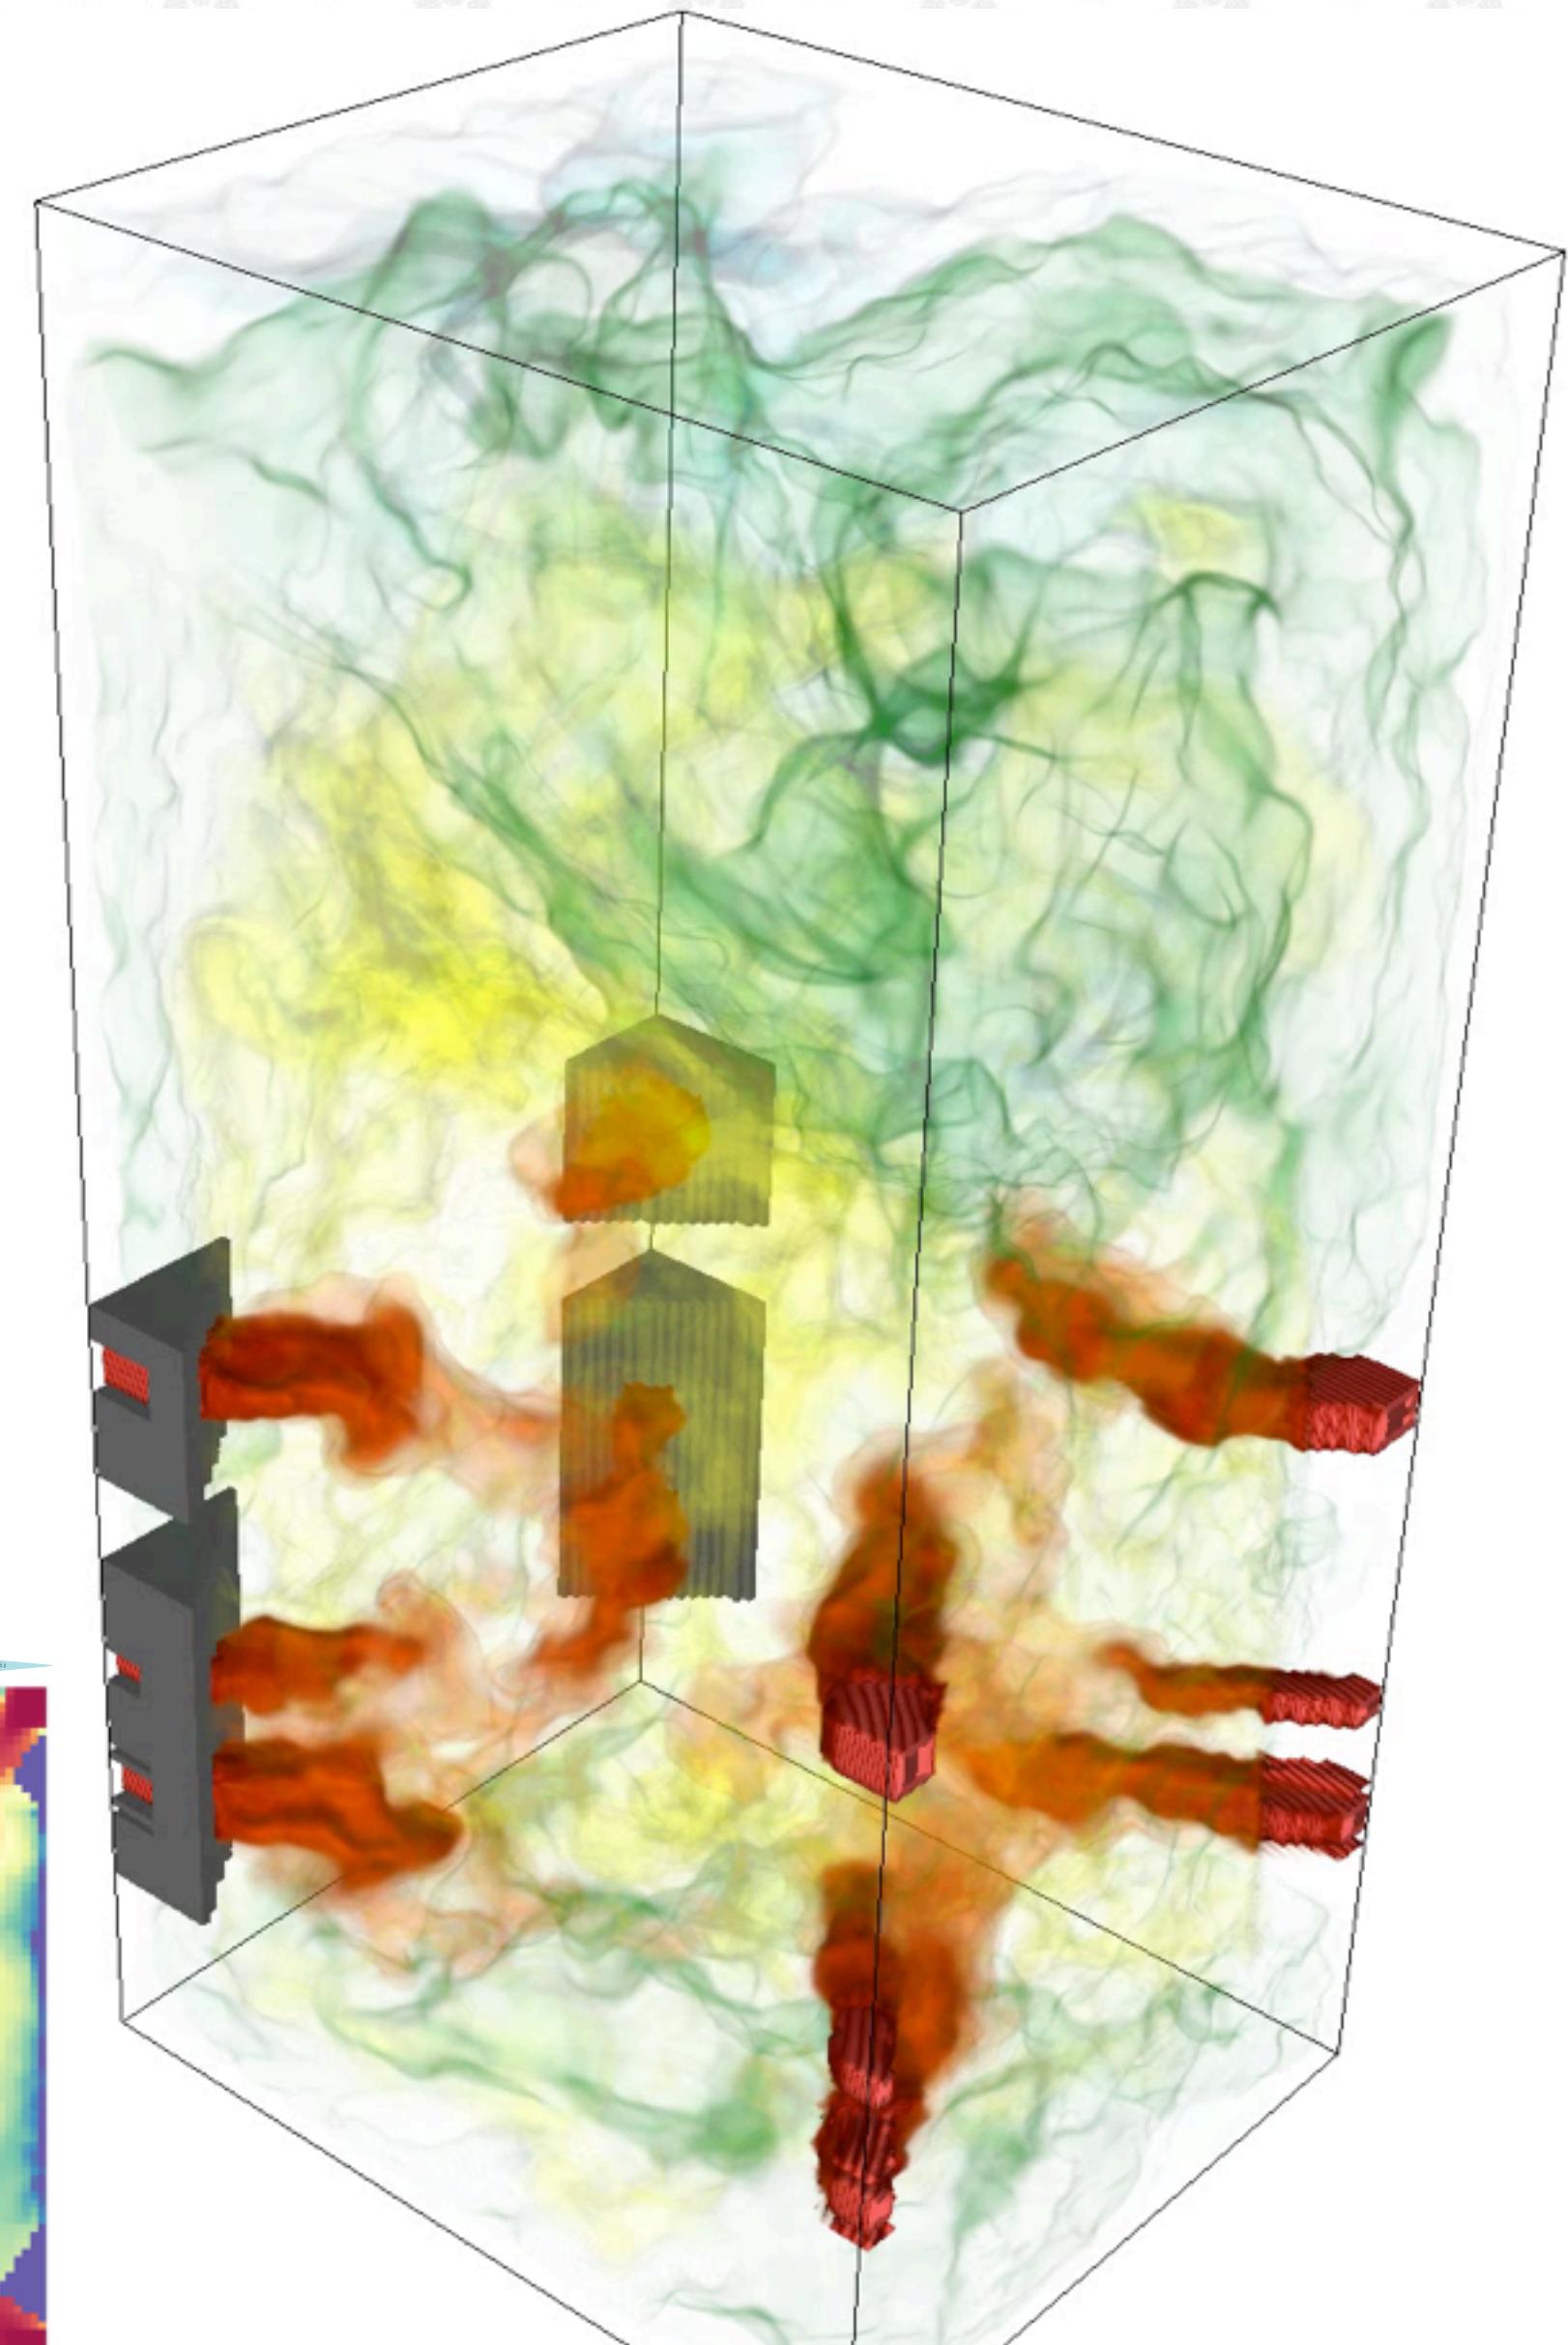
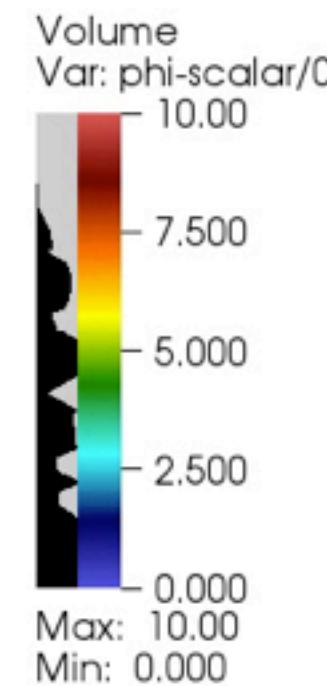
```
phi <= divX( interpX(lambda) * gradX(temperature) )  
      + divY( interpY(lambda) * gradY(temperature) )  
      + divZ( interpZ(lambda) * gradZ(temperature) );
```

```
// field type inference:  
typedef FaceTypes<FieldT>::XFace XFluxT;  
typedef FaceTypes<FieldT>::YFace YFluxT;  
typedef FaceTypes<FieldT>::ZFace ZFluxT;  
  
// operator type inference:  
typedef OpTypes<FieldT>::DivX DivX;  
typedef OpTypes<FieldT>::DivY DivY;  
typedef OpTypes<FieldT>::DivZ DivZ;
```

- One inlined grid loop, no temporaries.
- Better performance & scalability than without chaining.
- Compile-time consistency checking (field-operator and field-field compatibility).
- Runtime consistency checks for ghost cell validity.

Performance on CFD problems

- Wasatch is 3-4x faster when using the EDSL in Uintah on the mini boiler problem.
 - Pure LES without scalar transport or other complicating factors.
 - Consistent with prior results showing up to 5x & 10x faster than sibling codes in Uintah on a Taylor-Green vortex problem.
- Scalable to 260K processes.
- This week, we introduced a new Nebo CPU backend that resulted in 20-30% performance gains immediately felt through all of Wasatch.



CARBON CAPTURE
MULTIDISCIPLINARY
SIMULATION CENTER

Multicore & GPU Performance

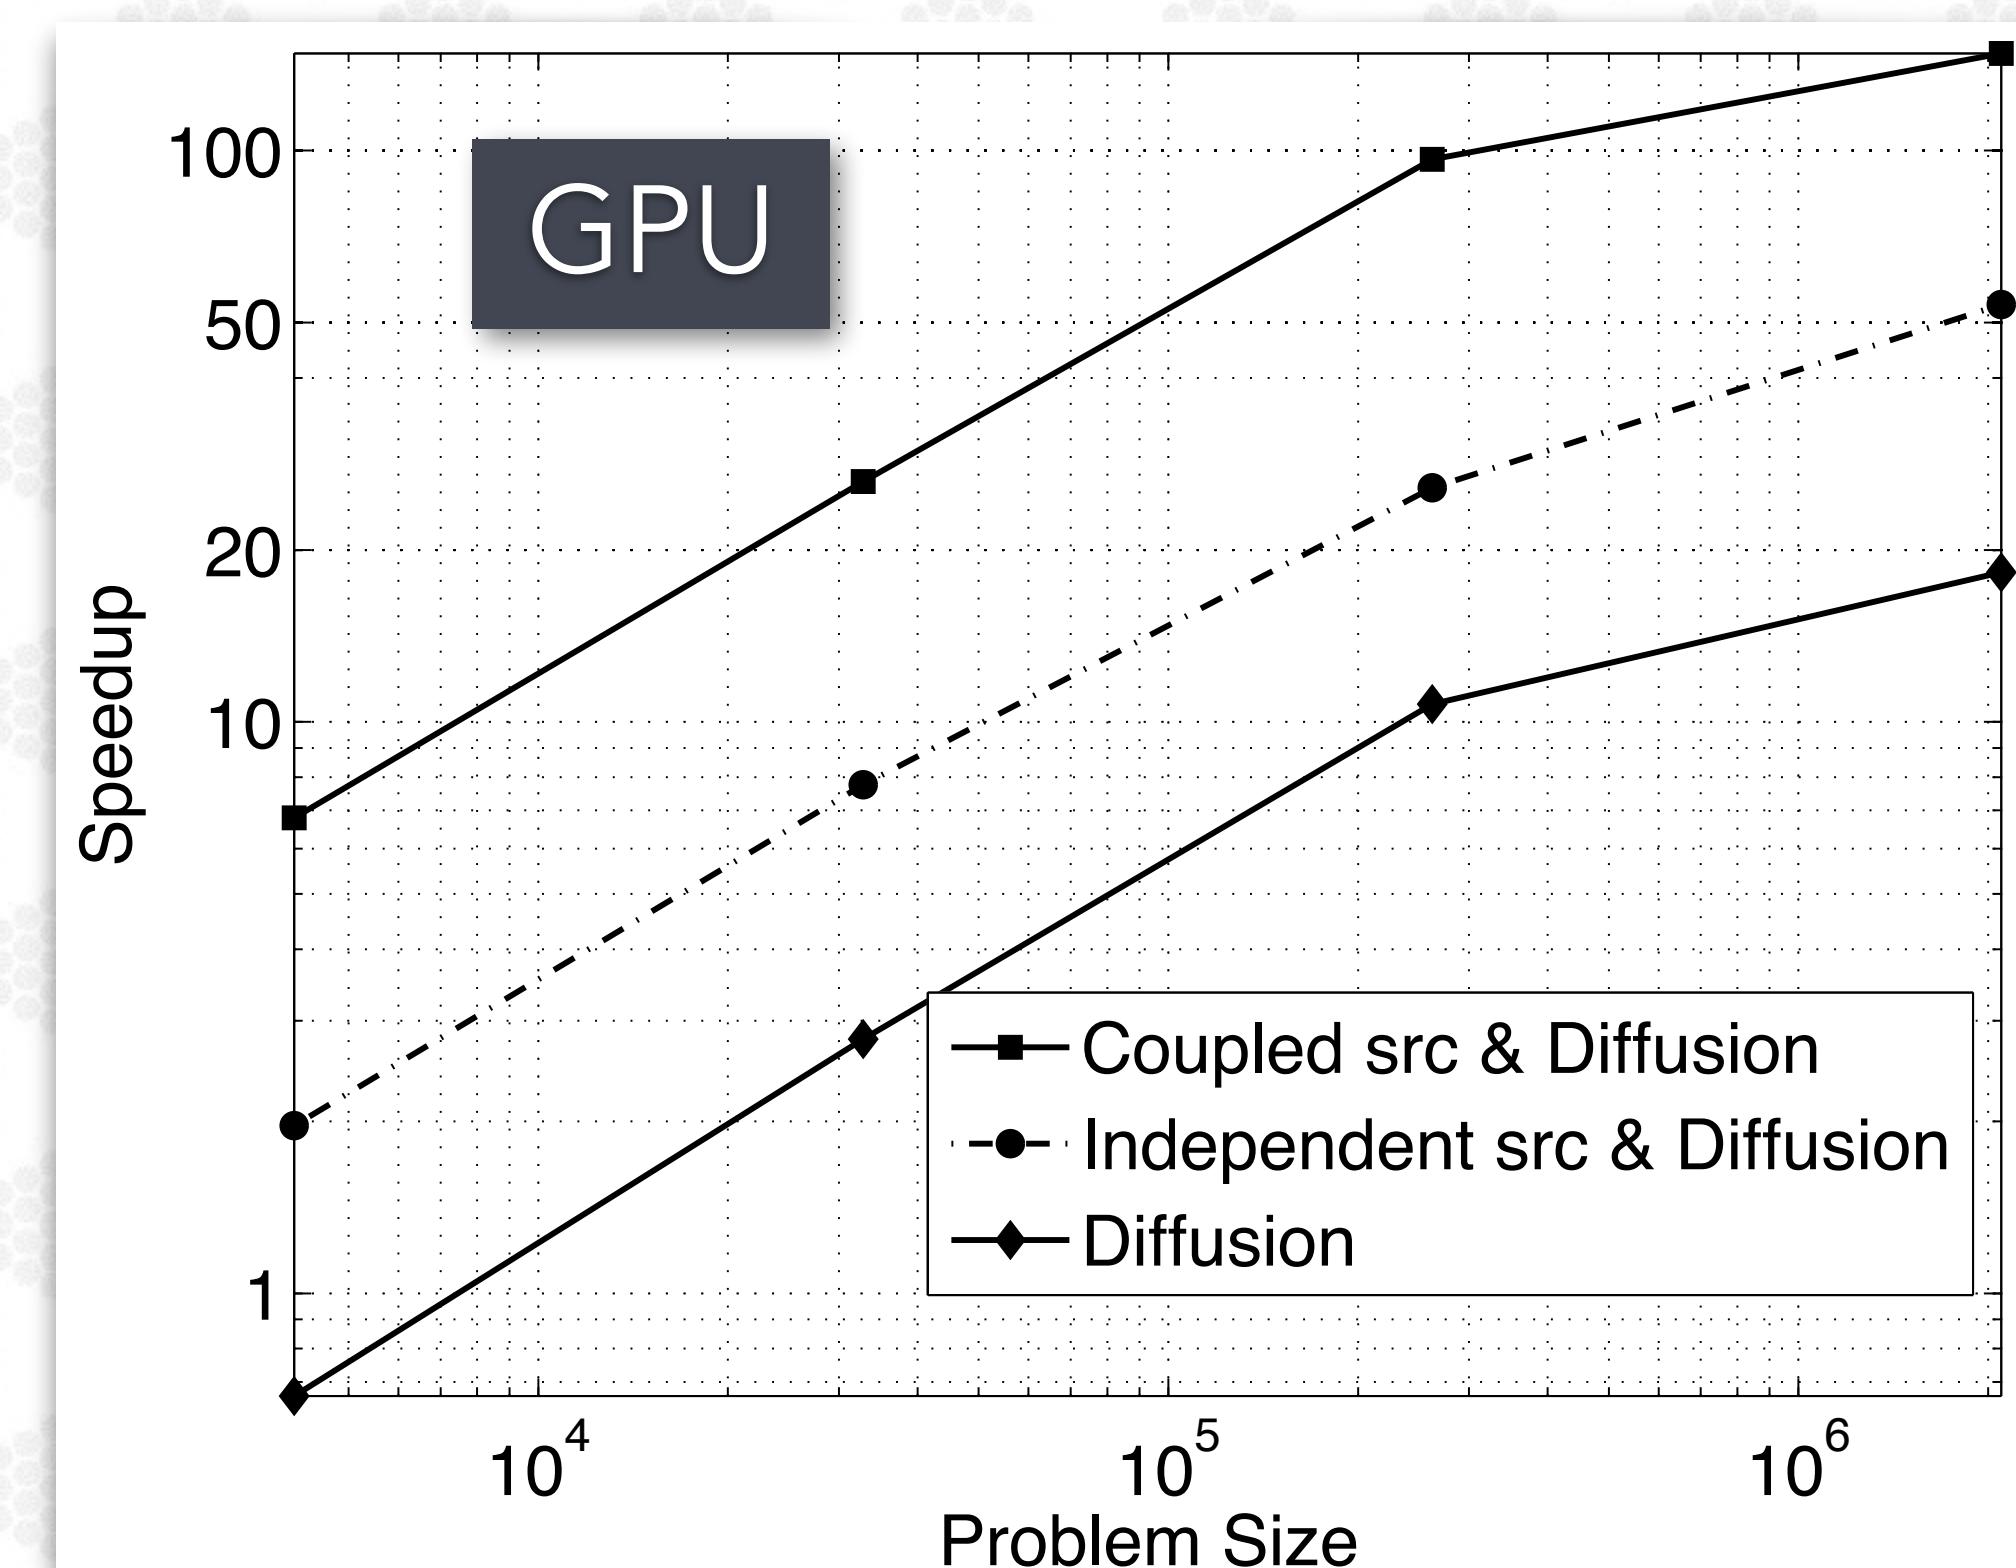
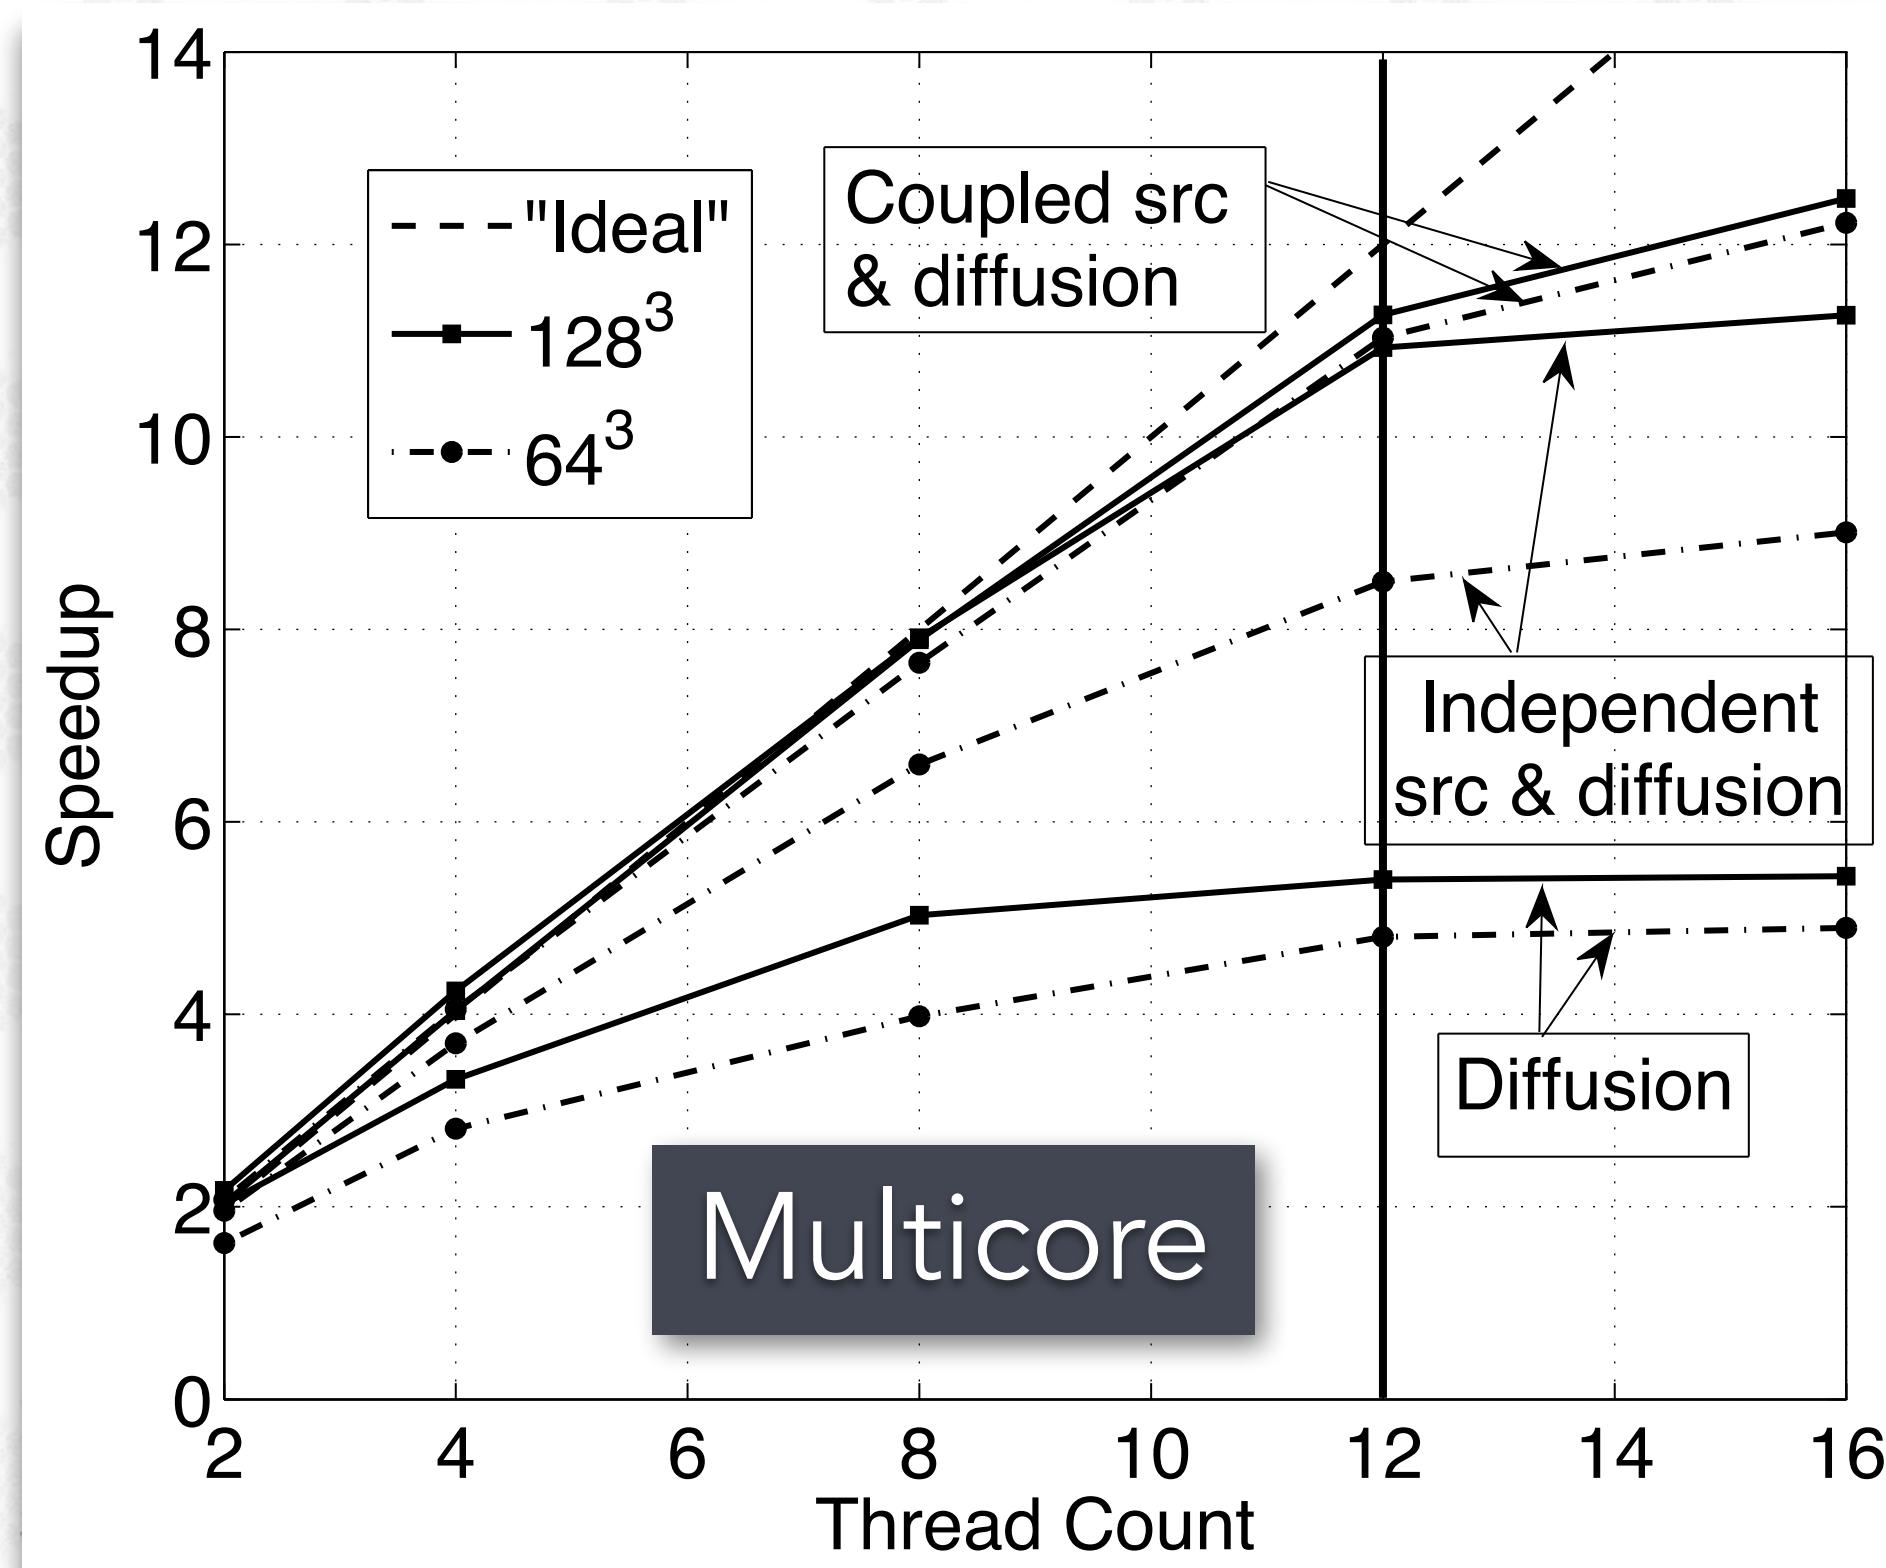
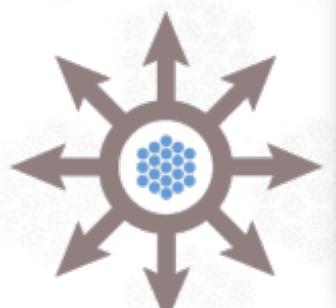
Test: mockup of a diffusion-reaction problem.

- Easily dial in the number of equations (30 here).
- Diffusion is an inexpensive stencil calculation.
- Reaction is an expensive point-wise calculation.

$$\frac{\partial \phi_i}{\partial t} = -\nabla \cdot \mathbf{J}_i + s_i$$

$$\mathbf{J}_i = -\Gamma \nabla \phi_i$$

$$s_i = f(\phi_i) \text{ or } s_i = f(\phi_j)$$



Parting Thoughts

- ➊ Hierarchical parallelization allows for flexible usage of available resources:

- *Domain decomposition (SIMD)*
 - Should allow a process to do computation on “interior” while waiting on communication from neighbors.
- *Task decomposition (MIMD)*
 - Decompose the solution into a DAG that can be scheduled asynchronously.
- *Vectorized parallel (SIMD)*
 - Break grid operations across multicore, GPU, etc.

Wasatch

ExprLib

SpatialOps

Uintah-X

- ➋ DAG representation is a scalable abstraction that:

- Handles problem complexity gracefully.
- Provides convenient separation of the problem’s structure from the data.
- Allows sophisticated scheduling algorithms to optimize scalability & performance.

- ➌ (E)DSLs are very useful

- Future-proofing: separate intent from implementation.
- **EDSLs** allow seamless transition of a code base and leverage existing compilers.
- Template metaprogramming pushes work from run-time to compile-time for more efficiency.

