



Flexible, Efficient Abstractions for High Performance Computation on Current and Emerging Architectures

JAMES C. SUTHERLAND

Associate Professor - Chemical Engineering

MATT MIGHT

Assistant Professor - School of Computing

CHRISTOPHER EARL

Postdoctoral Researcher

TONY SAAD

Research Associate

ABISHEK BAGUSETTY

M.S. Student







US DOE award
DE-NA0000740

NSF PetaApps
award 0904631



Motivation

-  Complex physics \Rightarrow complex software!
 - *is this necessary, or just a result of looking at the problem in the wrong way?*
-  Changing from model "A" to model "B" ...
 - *may require different transport equations*
 - *may introduce different nonlinear coupling*
-  Spatial discretization frequently permeates software design
 - *Model developers typically must deal with "mesh loops"*
 - often resort to "copy/paste/modify" tactics that are highly bug-prone
 - *Future proofing:*
 - What if you want to do OpenMP on these loops?
 - What happens when you learn that OpenMP is not the right tool?
 - pthreads, CUDA / OpenCL ... ?
-  Questions: Can we write efficient software that...
 - *... naturally handles complexity and allows us to easily extend/replace existing models?*
 - *... allows programmers to easily and robustly express intent while not worrying about "details?"*
 - *... allows us to refactor for different hardware architectures without rewriting the code base?*

Flexible...

Register all expressions

- Each "expression" calculates one or more field quantities.
- Each expression advertises its direct dependencies.

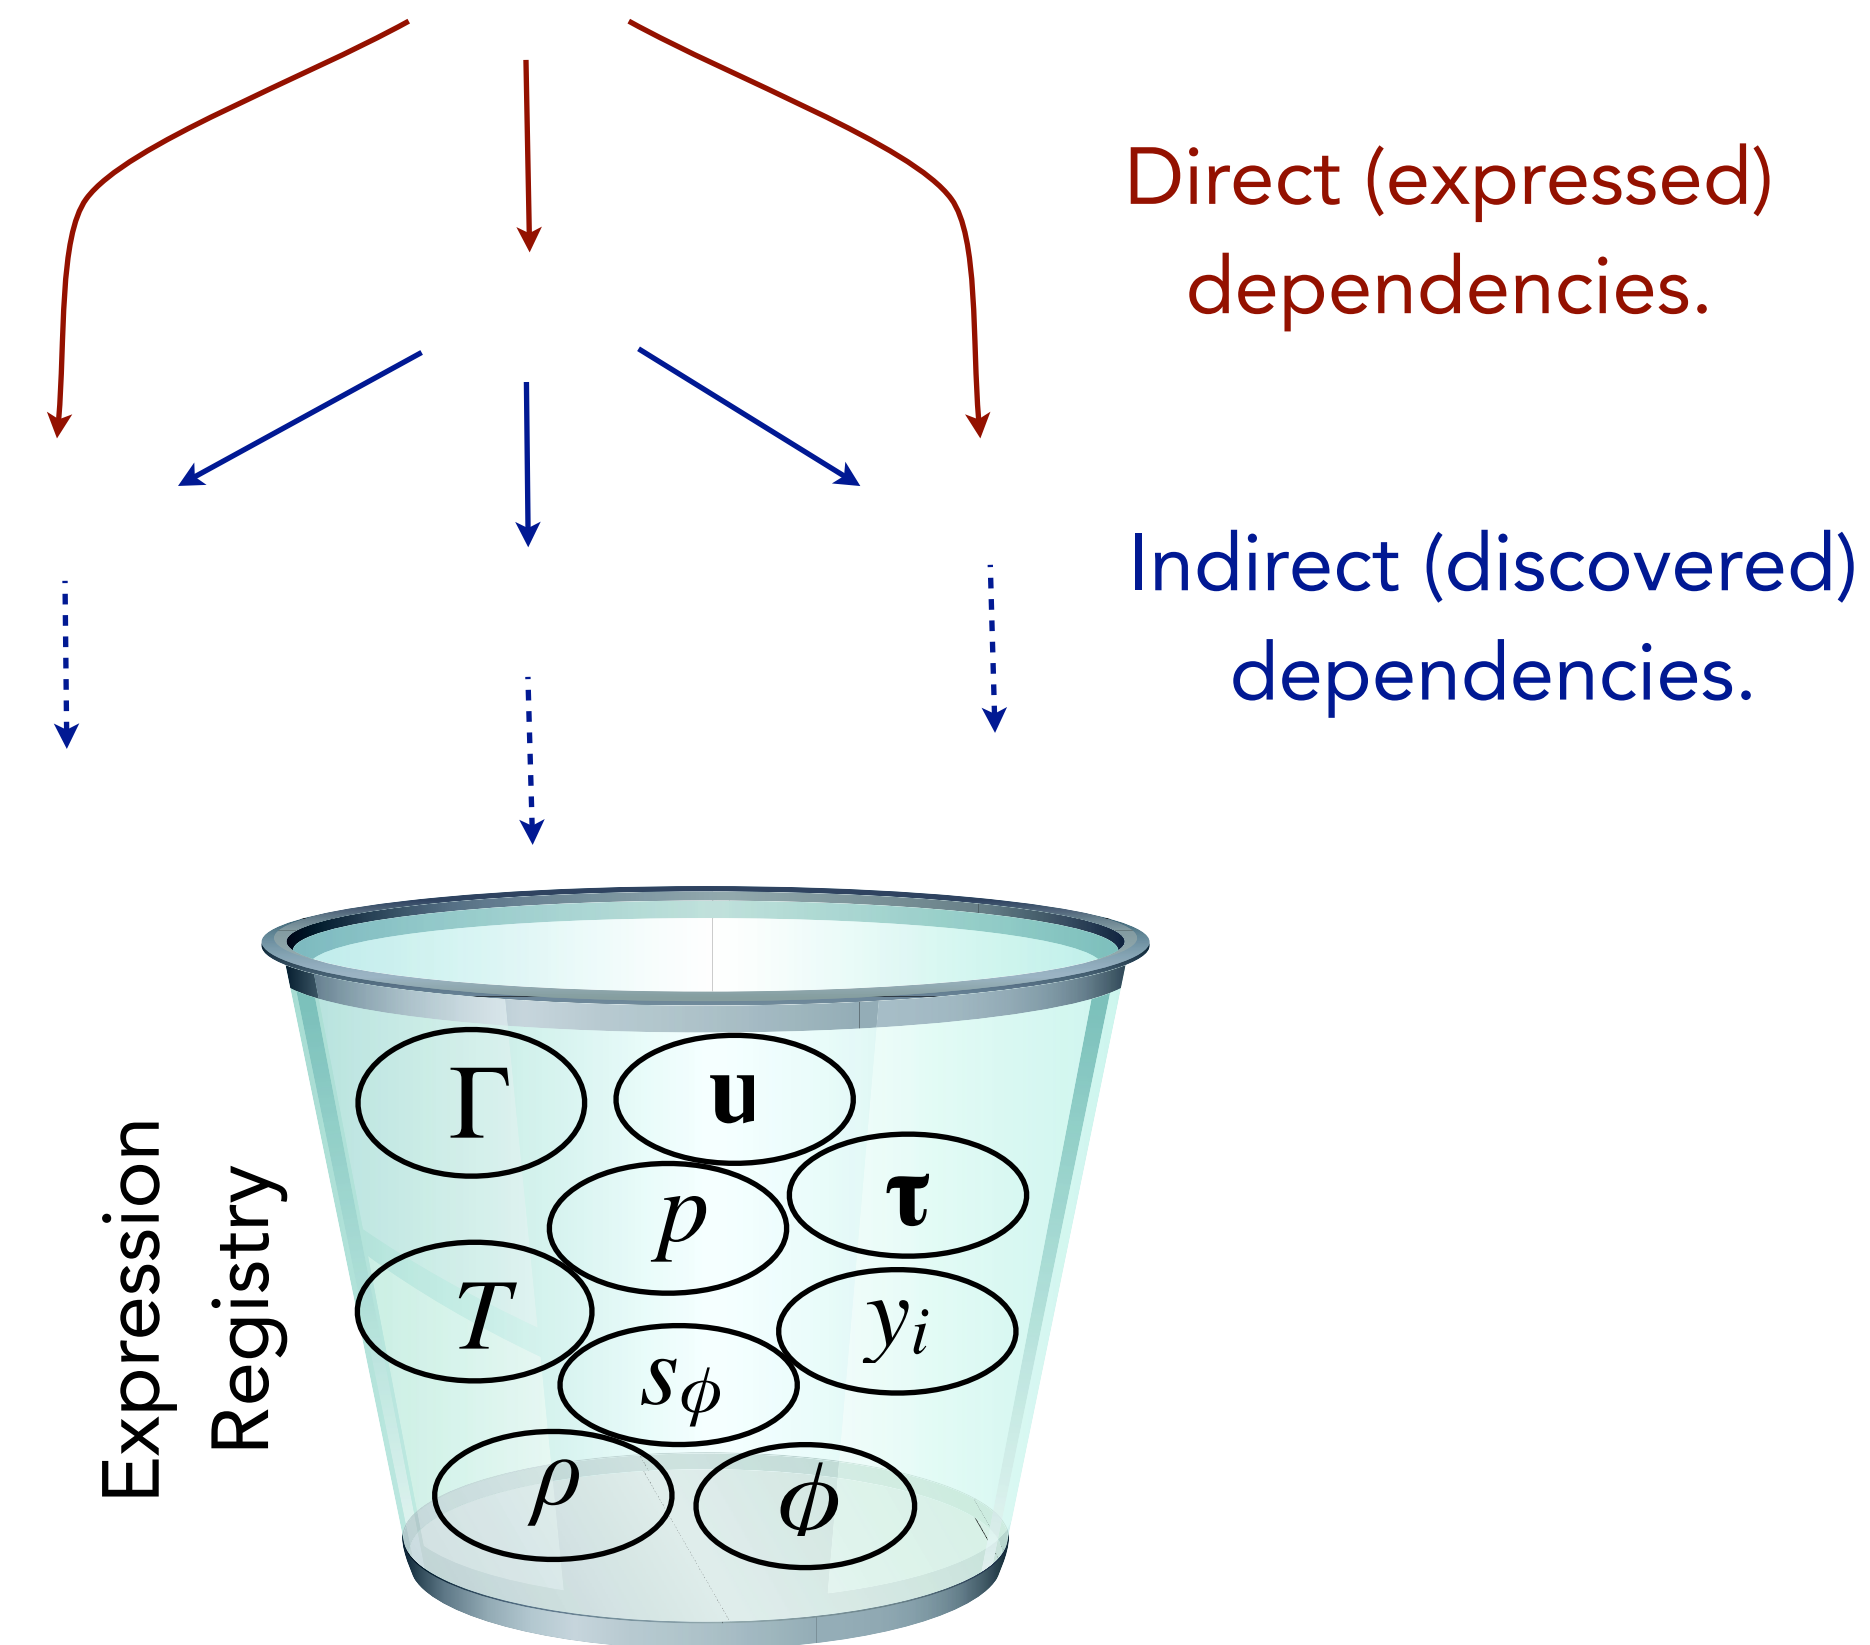
Set a "root" expression; construct a graph

- All dependencies are discovered/resolved automatically.
- Highly localized influence of changes in models.
- Not all expressions in the registry may be relevant/used.

From the graph:

- Deduce storage requirements & allocate memory (externally to each expression).
- Automatically schedule evaluation, ensuring proper ordering.
 - Asynchronous execution is critical! (overlap communication & computation)
- Robust scheduling algorithms are key.

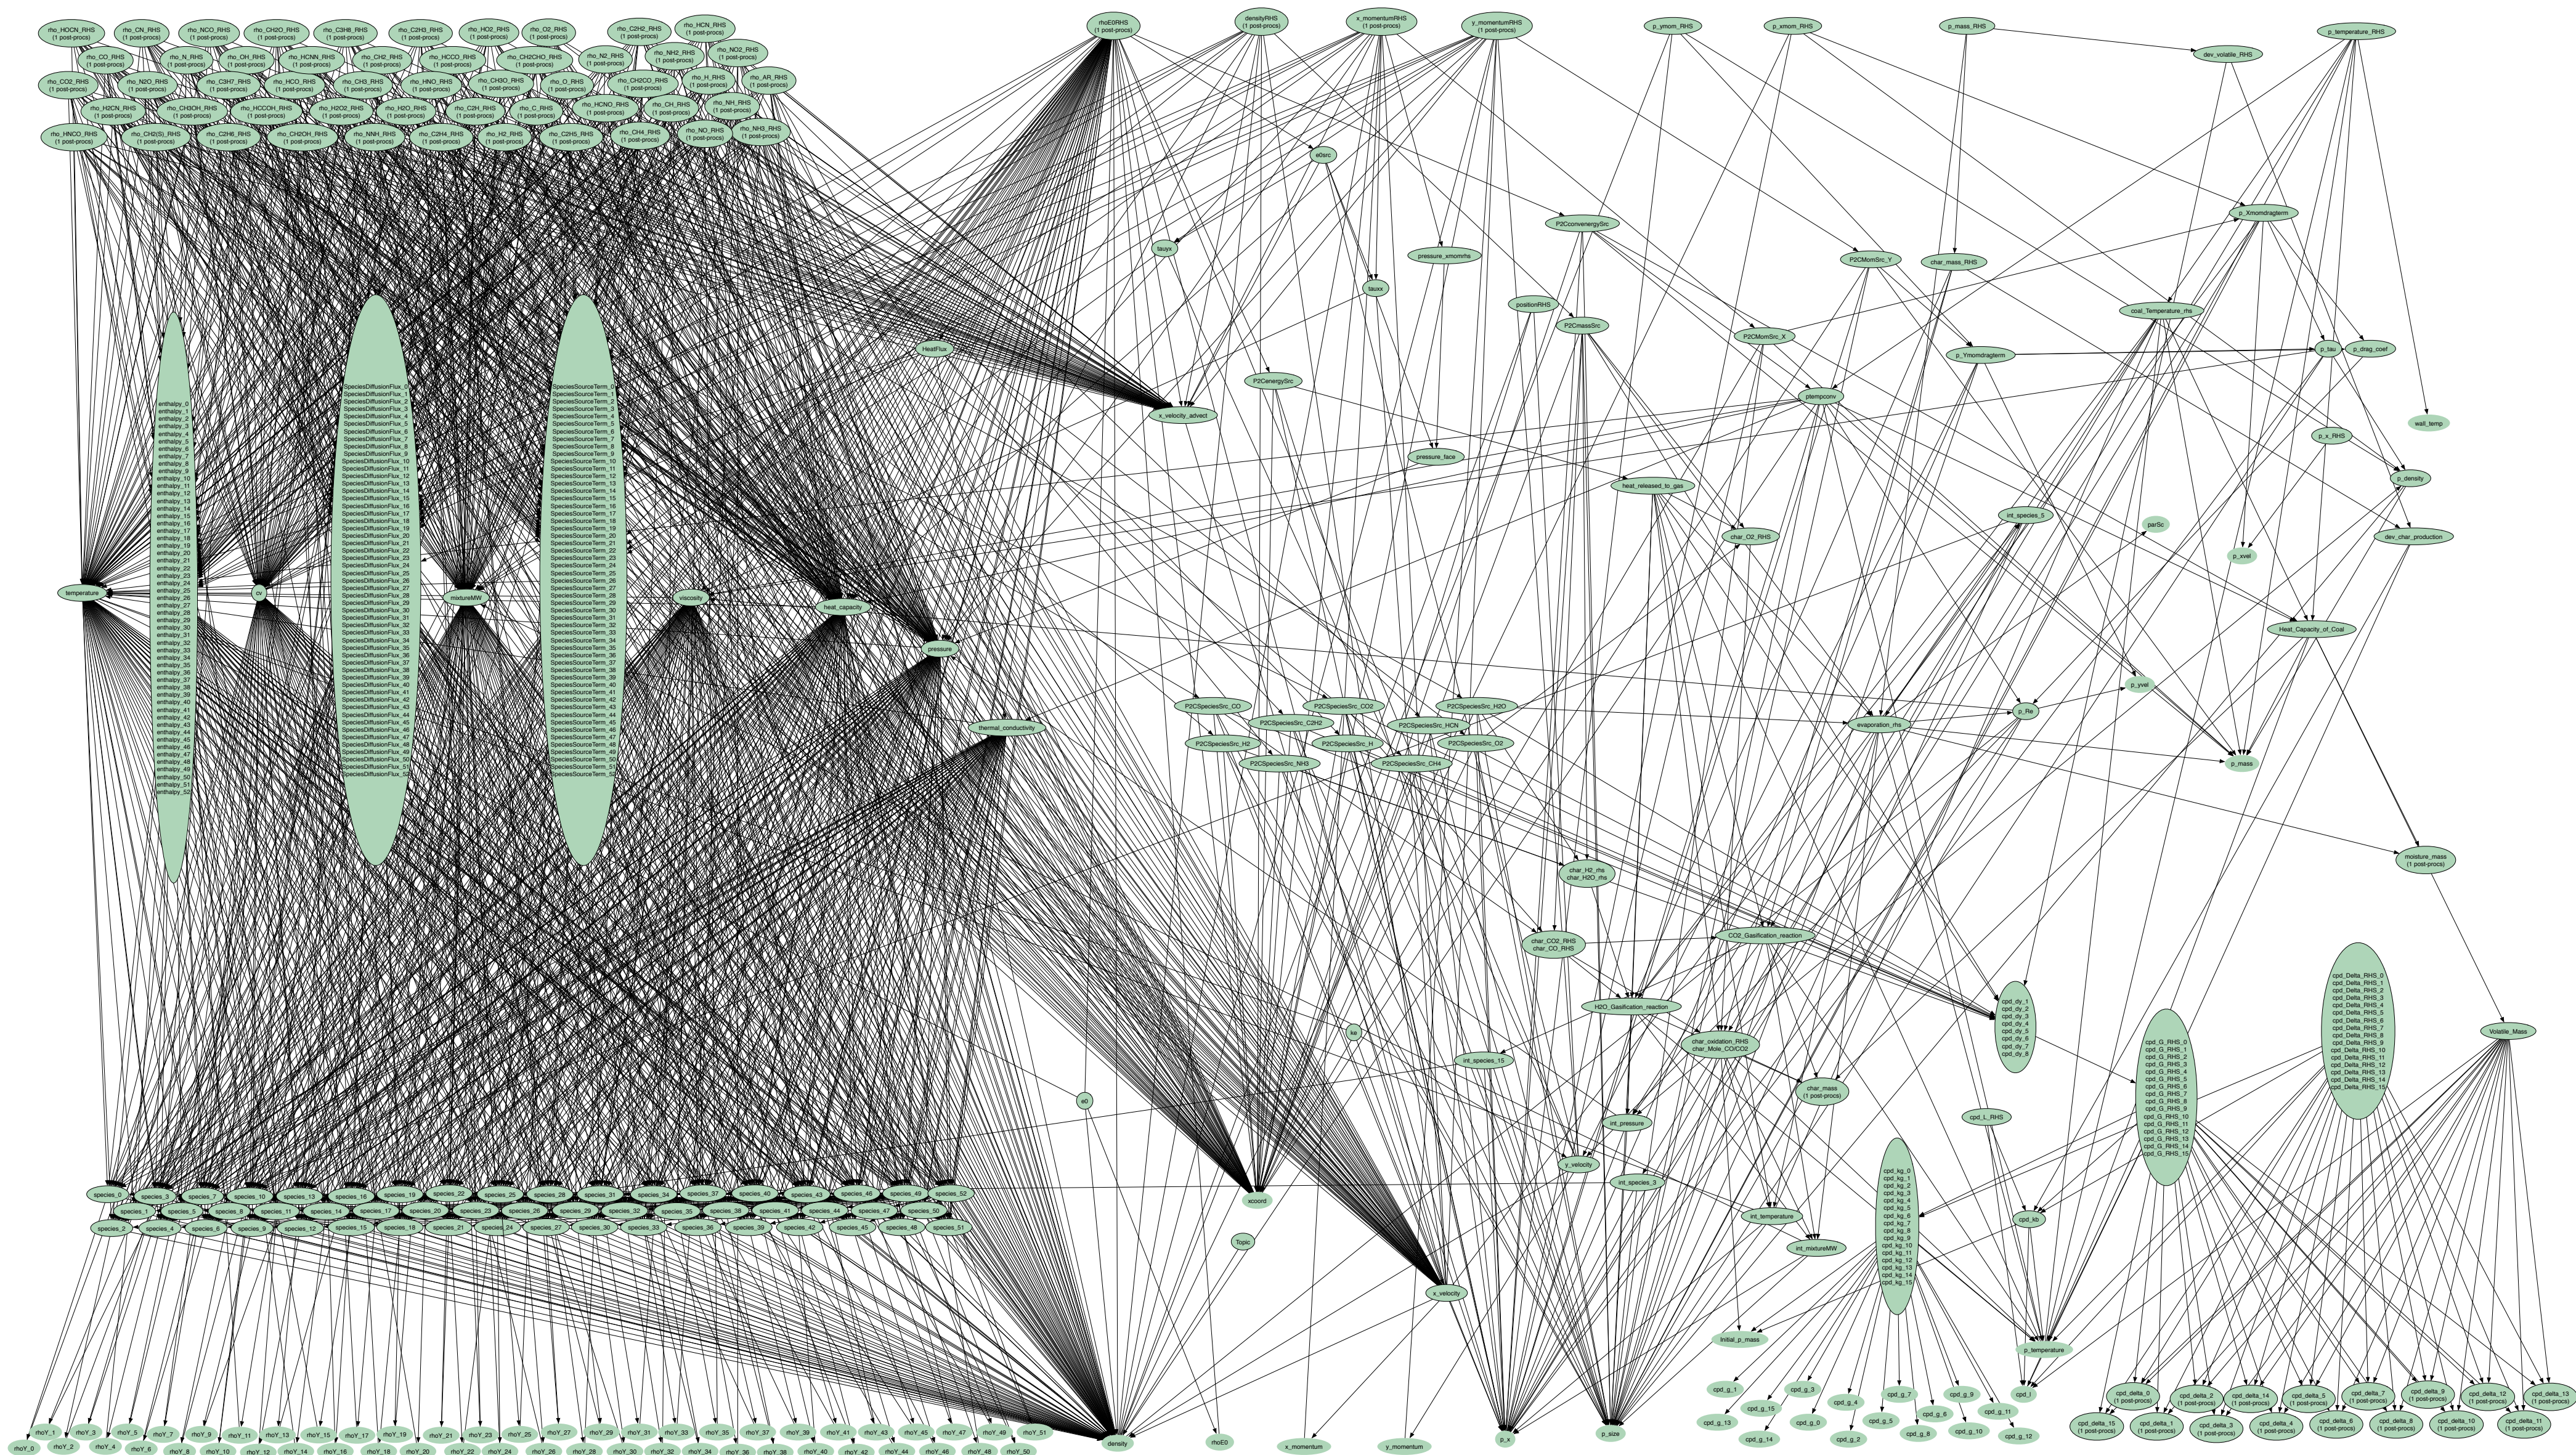
$$\Gamma = \Gamma(T, p, y_i)$$



*Notz, Pawlowski, & Sutherland (2012). ACM Transactions on Mathematical Software, 39(1).

Example: coal combustion

- 55 PDEs
- ~35 ODEs per particle
- Complex interphase coupling



Efficient...

Expressive syntax (matlab-style array operations)

- *Programmer expresses intent (problem structure) - not implementation.*

High performance

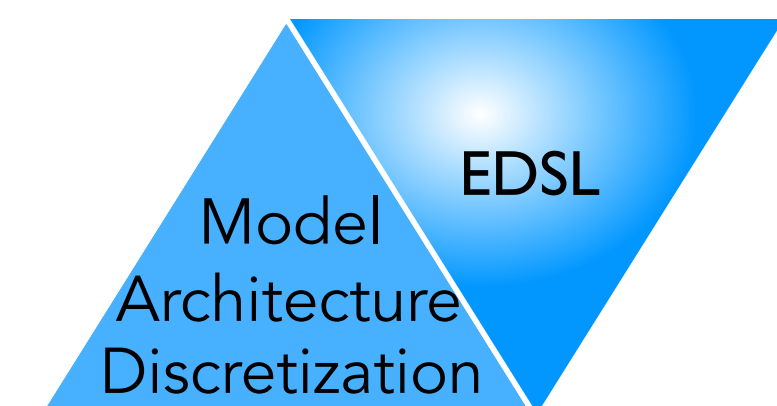
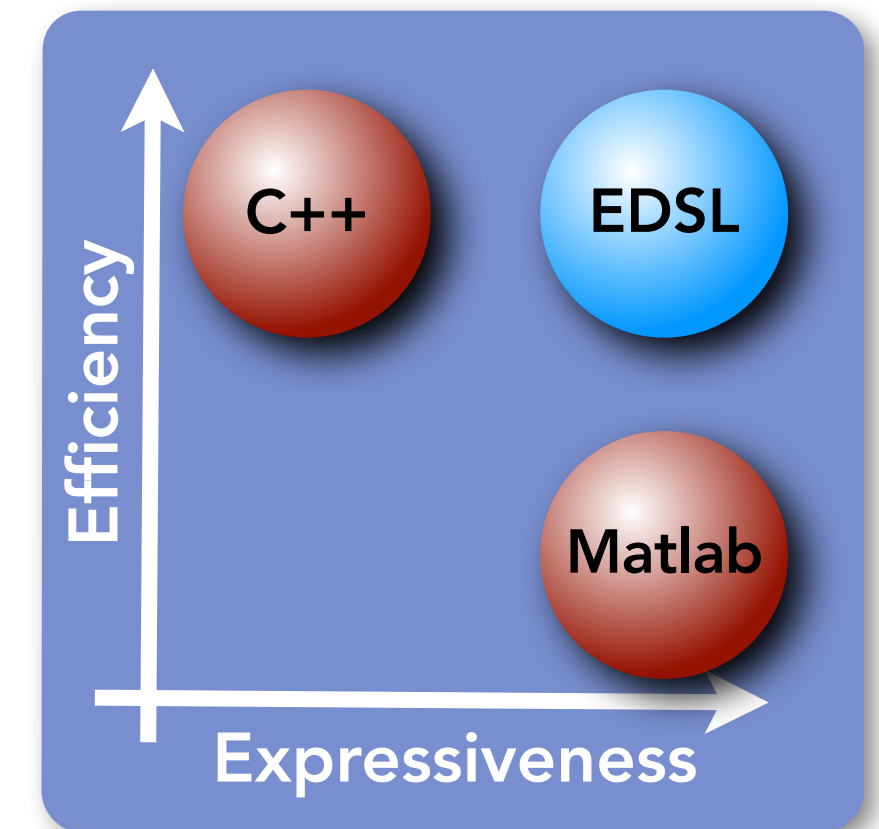
- *Should match hand-tuned code in performance.*

Extensible

- *Insulate programmer from architecture changes (e.g. multicore → GPU → ...).*
- *EDSL "back-end" compiles into code for target architecture.*

"Plays well with others"

- *Allow programmer to write in C++ and inter-operate with EDSL.*
- *Not an "all-or-none" approach: enable incremental adoption.*
- *Allows concurrent development of EDSL and application codes.*



Field Expressions

$$\vec{c} = \vec{a} + \sin(\vec{b})$$

Manual C++

Thread 1 {
 Field::const_iterator ia₁ = a₁.begin();
 Field::const_iterator ib₁ = b₁.begin();
 for(Field::iterator ic₁ = c₁.begin();
 ic₁ != c₁.end();
 ++ic₁, ++ia₁, ++ib₁) {
 *ic₁ = *ia₁ + sin(*ib₁);
 };
 :
Thread n {
 Field::const_iterator ia_n = a_n.begin();
 Field::const_iterator ib_n = b_n.begin();
 for(Field::iterator ic_n = c_n.begin();
 ic_n != c_n.end();
 ++ic_n, ++ia_n, ++ib_n) {
 *ic_n = *ia_n + sin(*ib_n);
 };

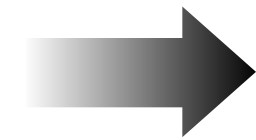
Nebo
EDSL

`c <<= a + sin(b);`

- Data parallel handled internally.
- Thread deployment (resizable threadpool).
- GPU deployment.
- Compile-time guarantee of field compatibility for given operations.

Chained Stencil Operations

$$\begin{aligned}\phi &= -\nabla \cdot \mathbf{q} \\ &= \nabla \cdot (\lambda \nabla T)\end{aligned}$$



```
phi <<= divX( interpX(lambda) * gradX(temperature) )
          + divY( interpY(lambda) * gradY(temperature) )
          + divZ( interpZ(lambda) * gradZ(temperature) );
```

// field type inference:

```
typedef FaceTypes<FieldT>::XFace  XFluxT;
typedef FaceTypes<FieldT>::YFace  YFluxT;
typedef FaceTypes<FieldT>::ZFace  ZFluxT;
```

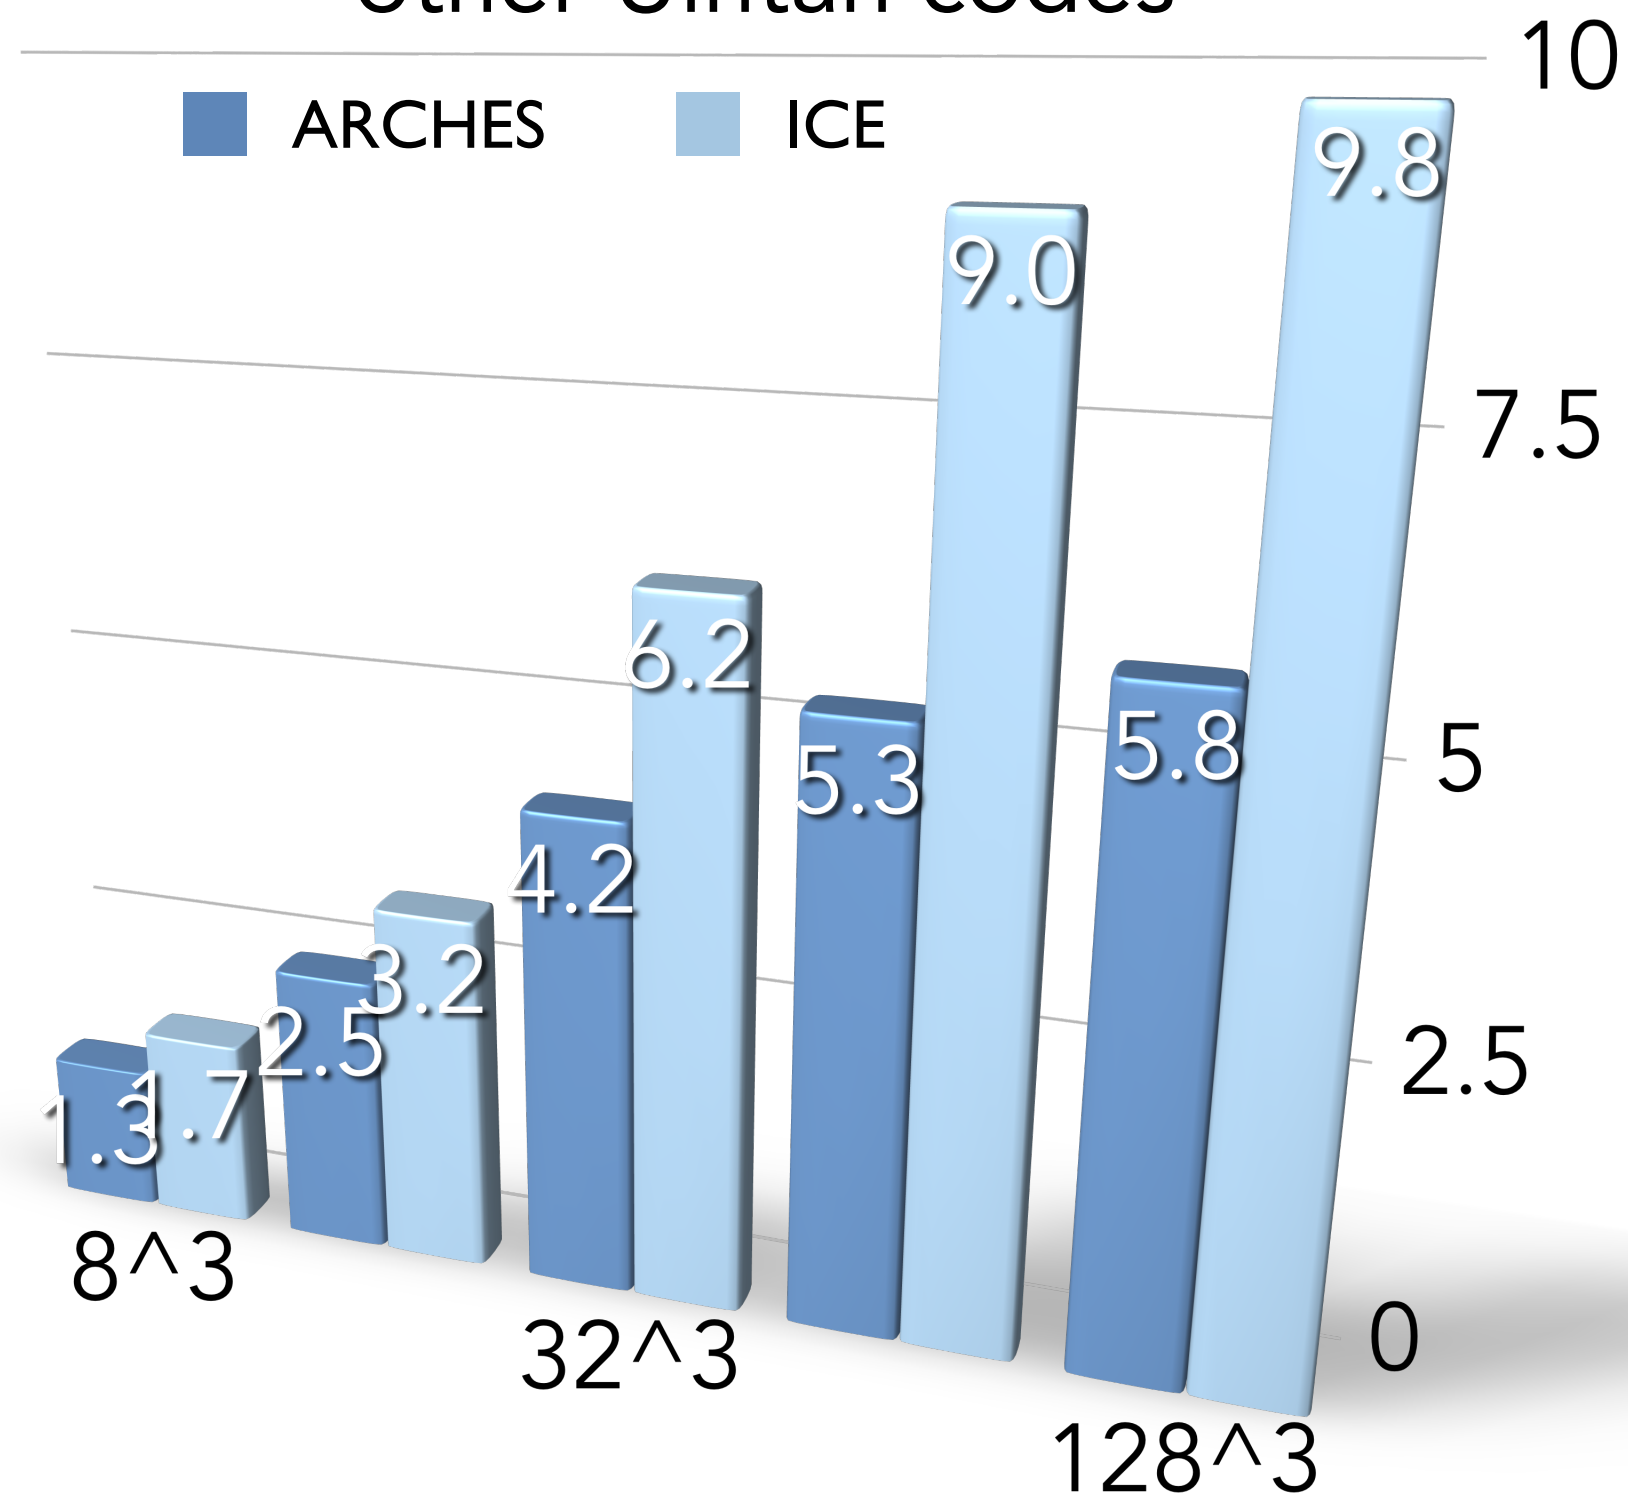
// operator type inference:

```
typedef OpTypes<FieldT>::DivX  DivX;
typedef OpTypes<FieldT>::DivY  DivY;
typedef OpTypes<FieldT>::DivZ  DivZ;
```

- One inlined grid loop, no temporaries.
- Better performance & scalability than without chaining.
- Compile-time consistency checking (field-operator and field-field compatibility).
- Runtime consistency checks for ghost cell validity.

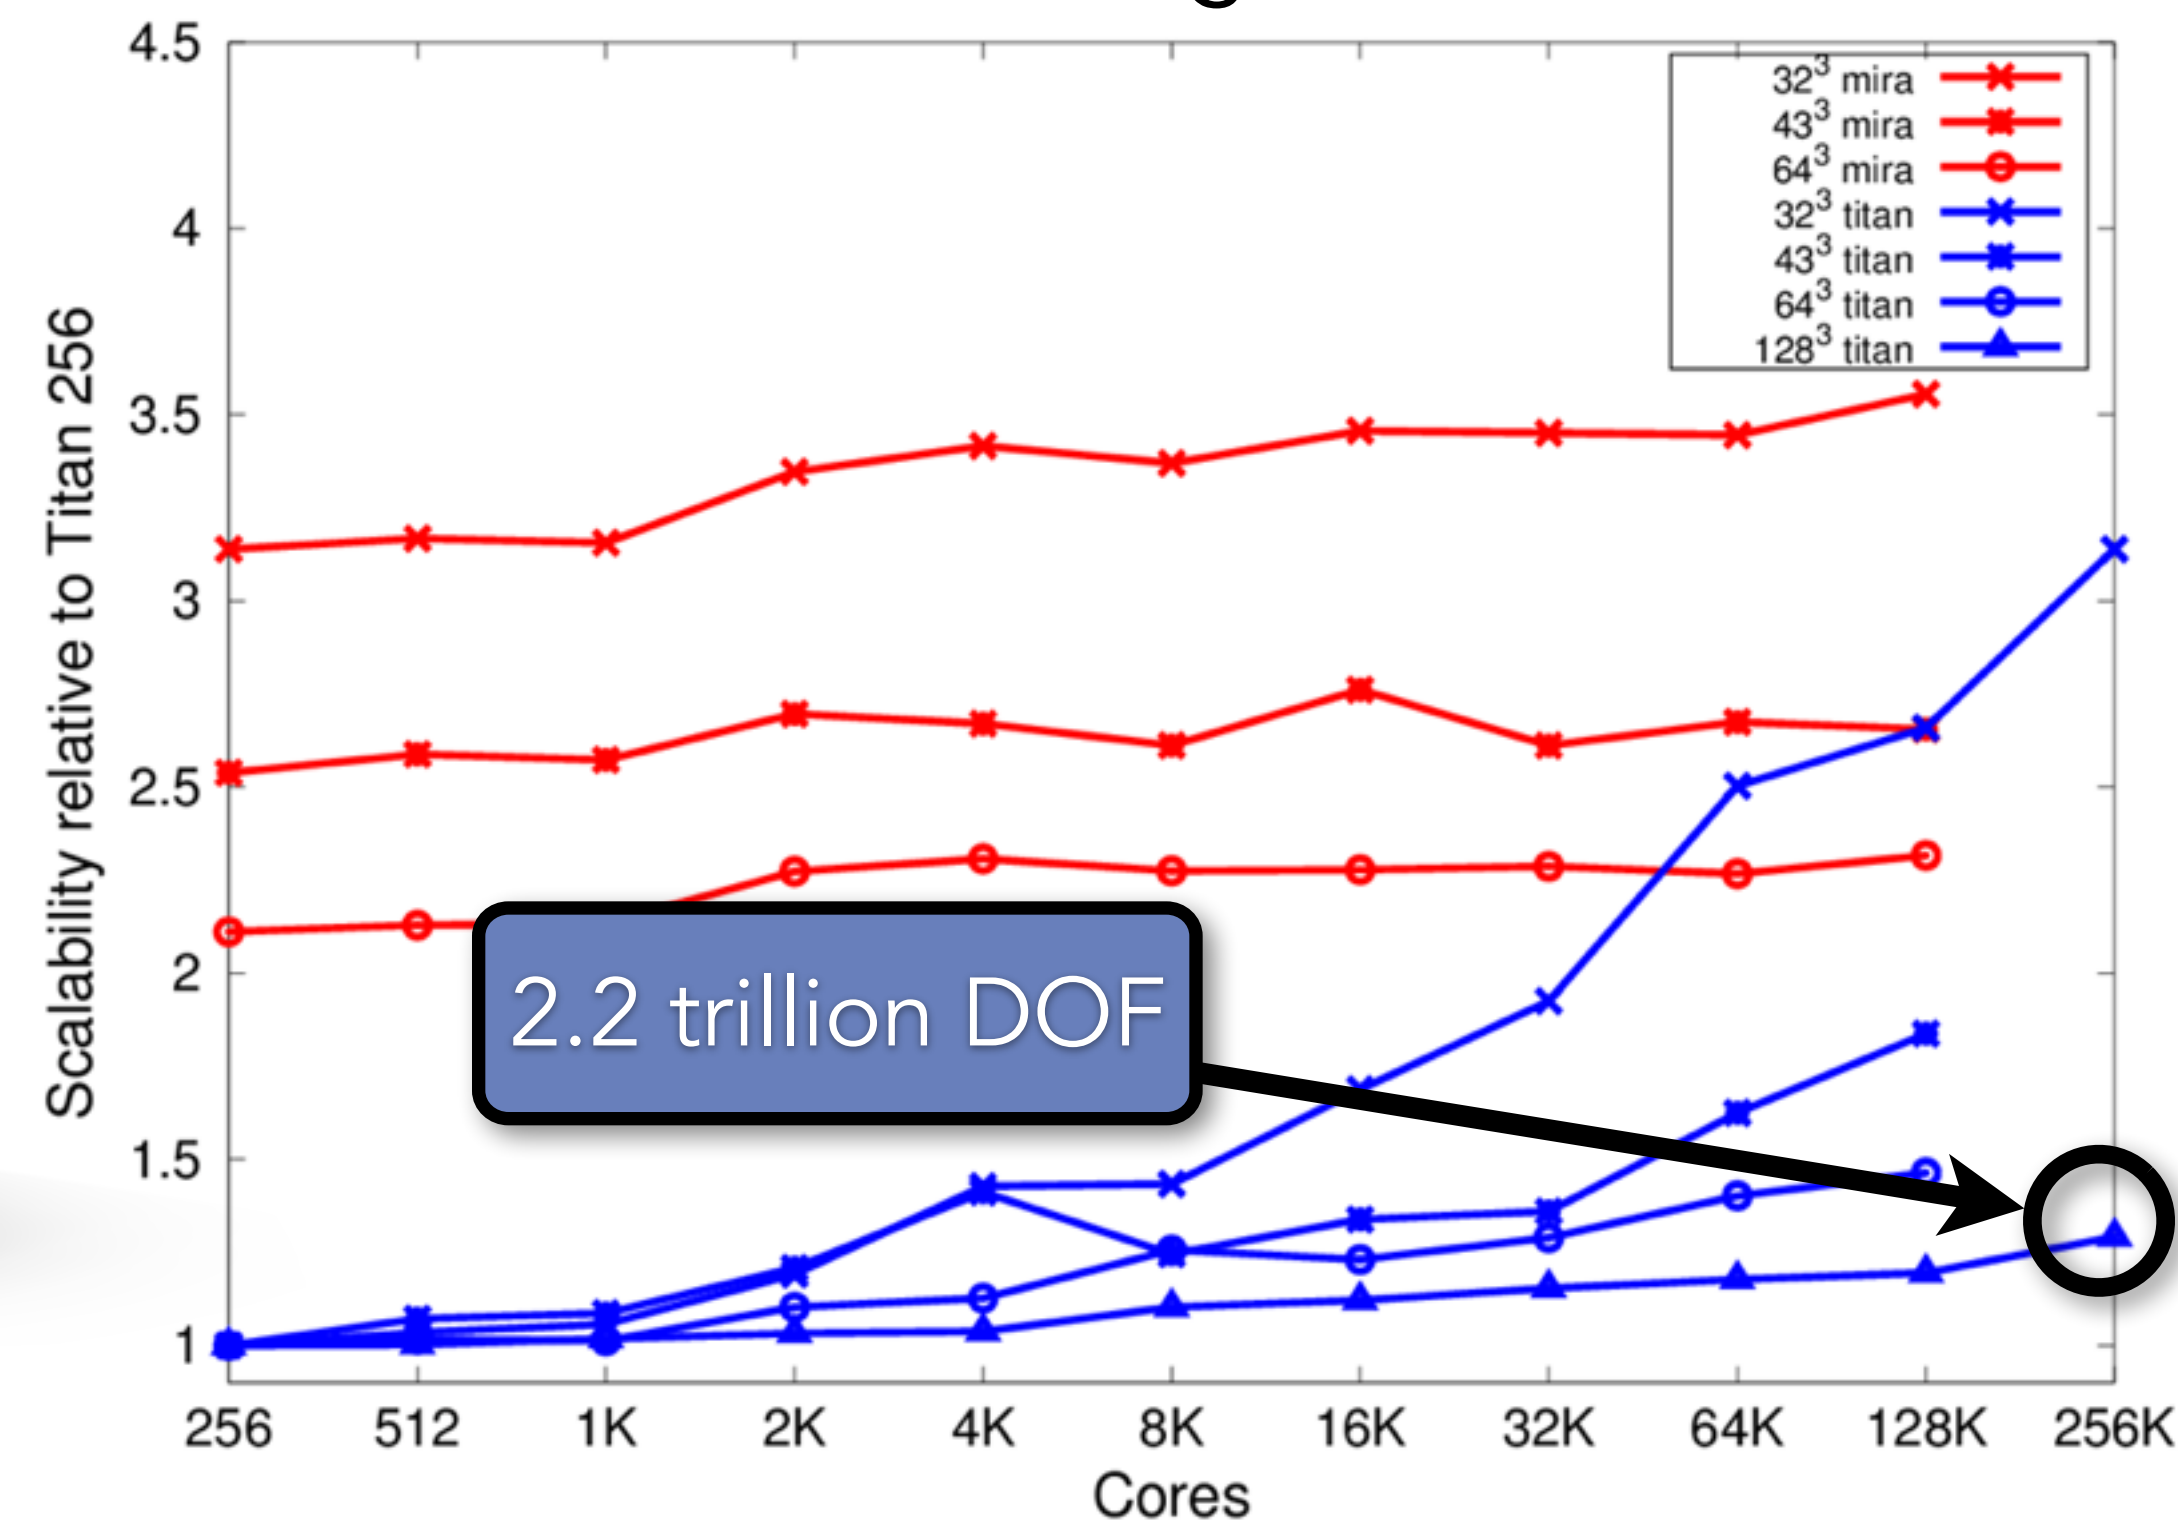
Putting it Together: Performance & Scalability

Speedup using DSL* relative to
other Uintah codes



*Comparison to ICE and ARCHES, sister codes in Uintah, on a 3D Taylor-Green vortex problem.
Run on a single processor.

Weak scaling on Titan



"1" indicates perfect weak scaling

Multicore & GPU Performance

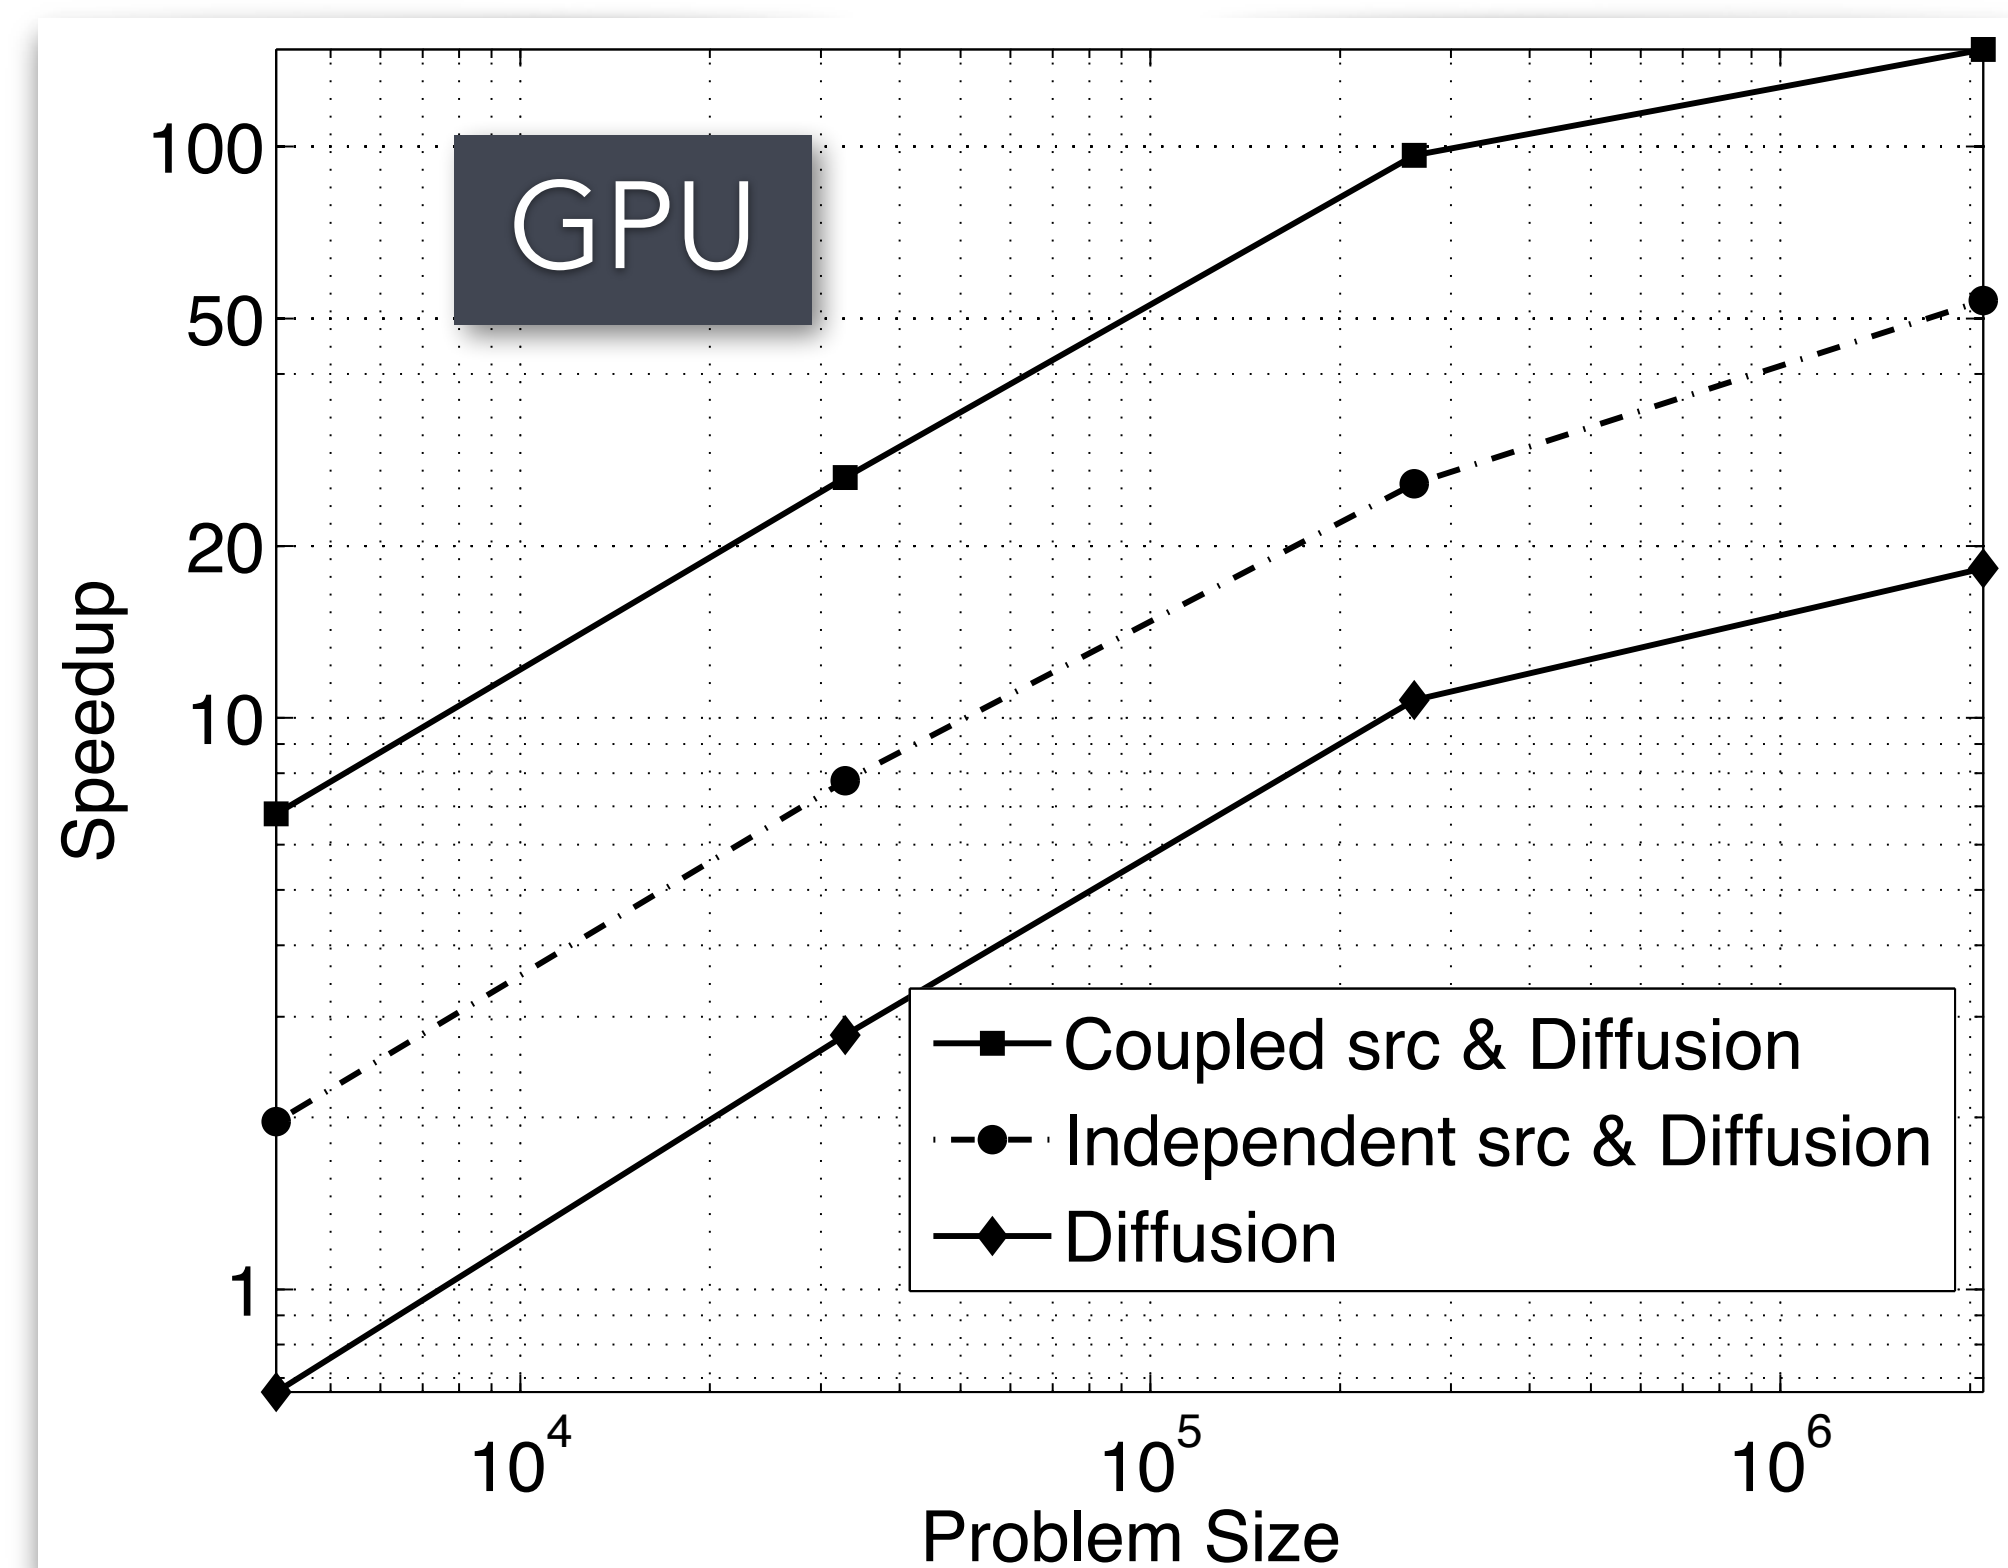
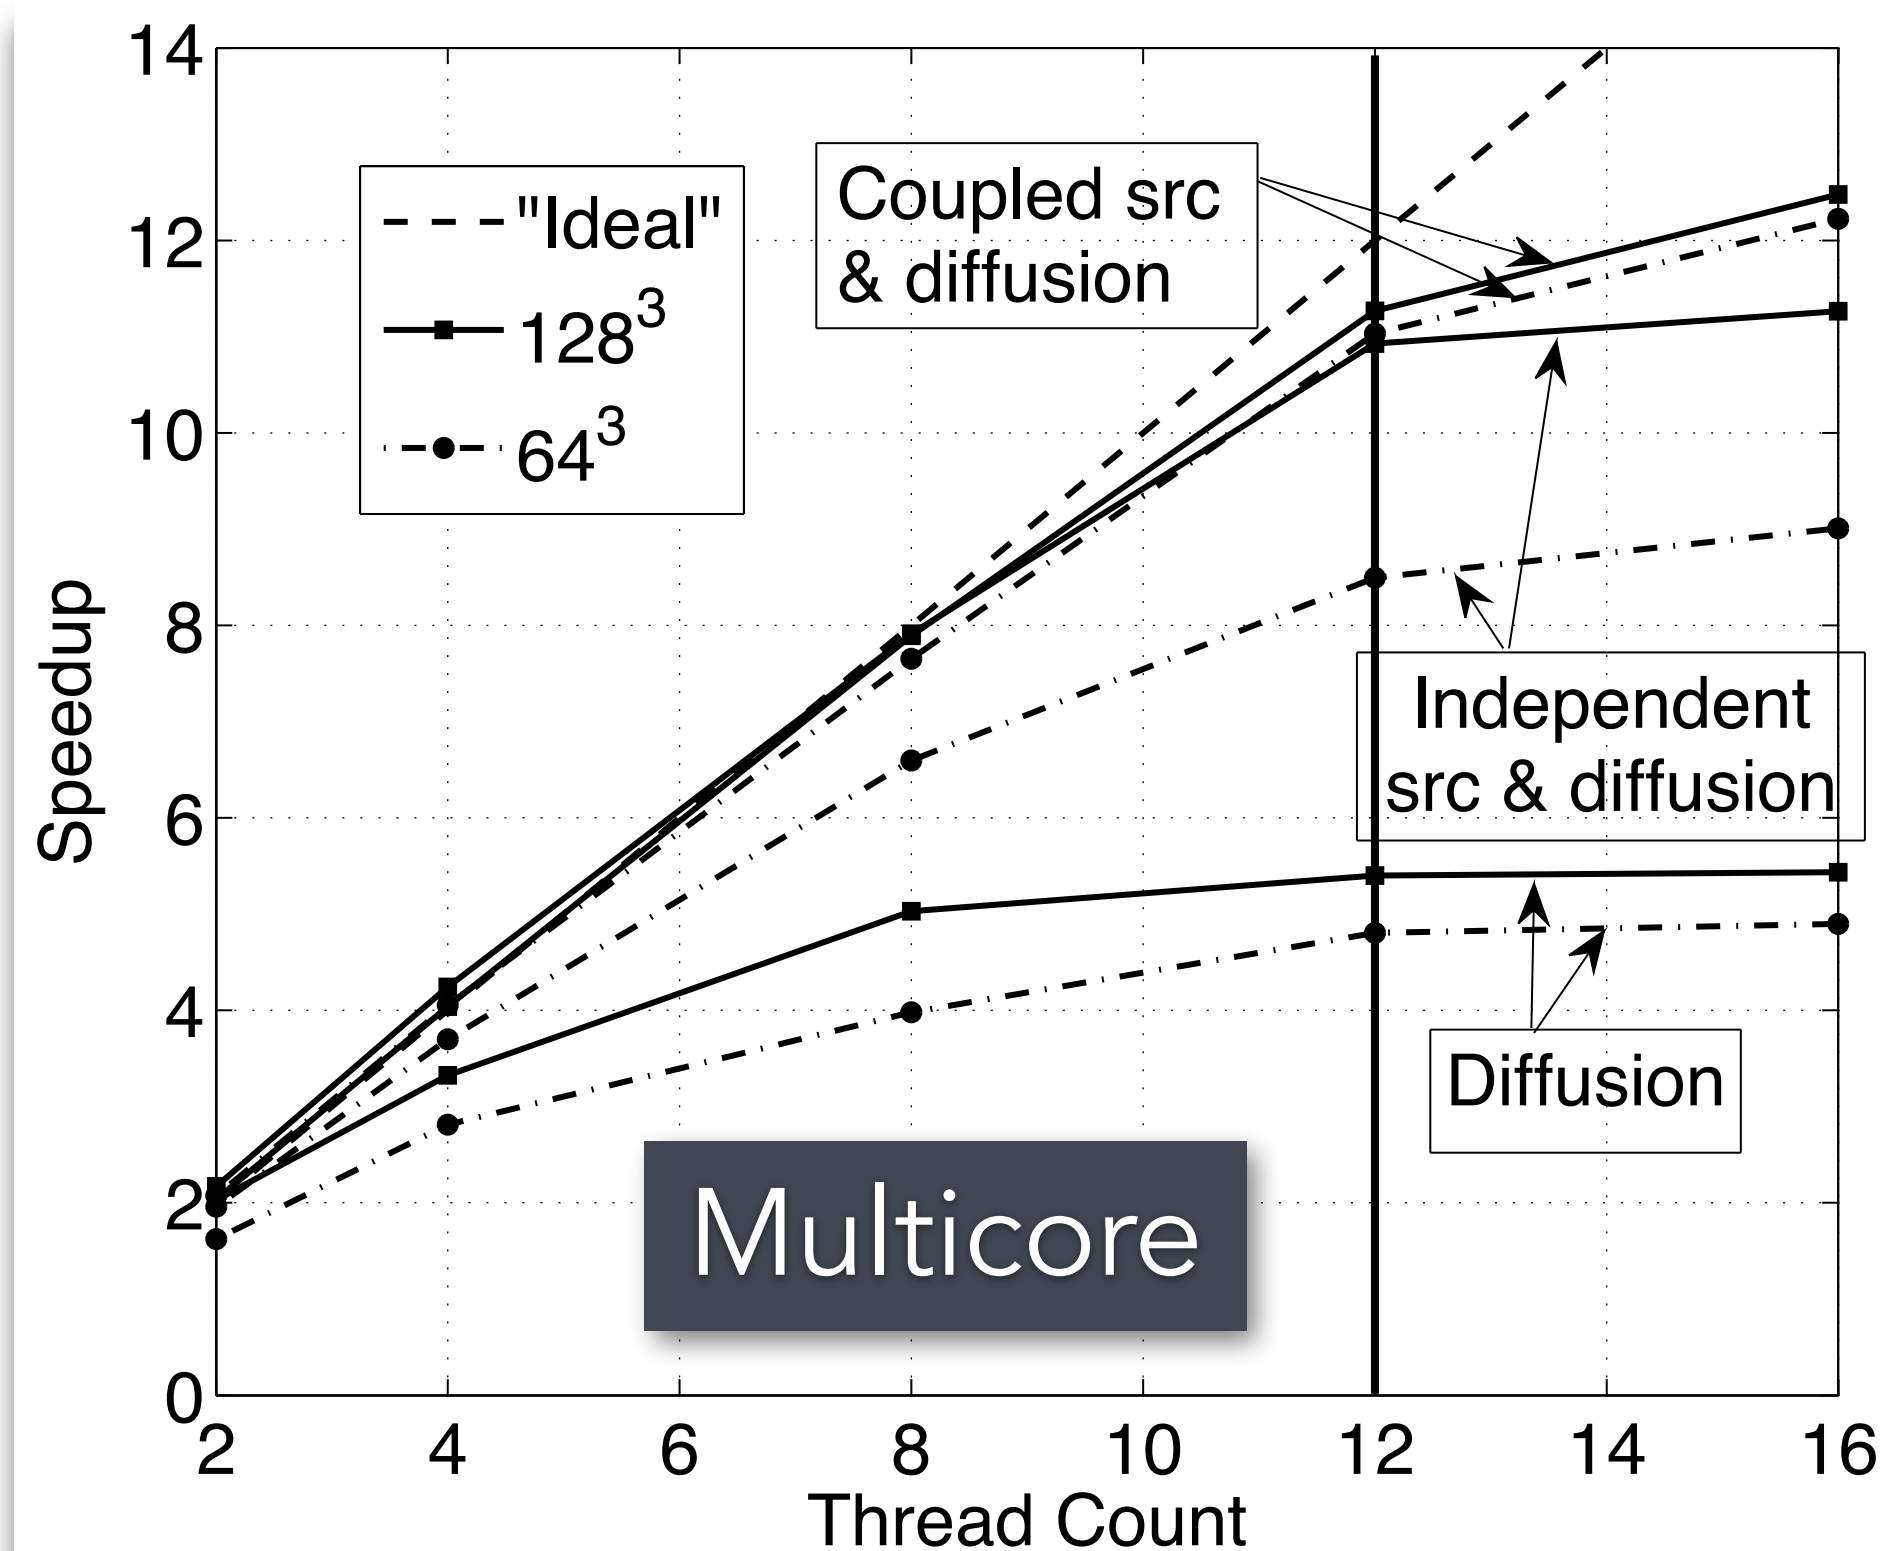
Test: mockup of a diffusion-reaction problem.

- Easily dial in the number of equations (30 here).
- Diffusion is an inexpensive stencil calculation.
- Reaction is an expensive point-wise calculation.

$$\frac{\partial \phi_i}{\partial t} = -\nabla \cdot \mathbf{J}_i + s_i$$

$$\mathbf{J}_i = -\Gamma \nabla \phi_i$$

$$s_i = f(\phi_i) \text{ or } s_i = f(\phi_j)$$



Parting Thoughts

 Hierarchical parallelization allows for flexible usage of available resources:

- *Domain decomposition (SIMD)*
 - Should allow a process to do computation on “interior” while waiting on communication from neighbors.
- *Task decomposition (MIMD)*
 - Decompose the solution into a DAG that can be scheduled asynchronously.
- *Vectorized parallel (SIMD)*
 - Break grid operations across multicore, GPU, etc.

 DAG representation is a scalable abstraction that:

- *Handles problem complexity gracefully.*
- *Provides convenient separation of the problem’s structure from the data.*
- *Allows sophisticated scheduling algorithms to optimize scalability & performance.*

 (E)DSLs are very useful

- *Future-proofing: separate intent from implementation.*
- *EDSLs allow seamless transition of a code base and leverage existing compilers.*
- *Template metaprogramming pushes work from run-time to compile-time for more efficiency.*